

ARTIFICIAL INTELLIGENCE SOFTWARE DEVELOPMENT

Week 10 Lecture 1
Dr. Hari M Koduvely

Agenda for Today

- ❑ Theory:
 - Fundamentals of Data Engineering :
 - Breakout session on Data Model & Data Pipeline for your project

Fundamentals of Data Engineering

Data Engineering refers to building systems for collecting, storing and processing large quantities of user and system generated data.

Data Sources

- ❑ An AI system may use data from different sources having different properties
- ❑ Understanding sources and properties of data can help in its efficient usage

Data Sources

□ User Input Data:

- Text, Videos, Uploaded files etc.
- Quite often not in the correct format.
 - E.g. Text instead of numbers, uploaded file with wrong extn etc.
- Needs more validation checking and processing.
- Often needs to process the data immediately and quickly

Data Sources

- System Generated Data:
 - System logs, model predictions etc.
 - Usually well formatted, hence requires less checking
 - Large volume of data
 - Important signals can be buried in the noise
 - Storage could be expensive
 - Typically processed in batches
 - System generated data can include user behavior
 - What user clicked, moved, sent etc.
 - Considered as part of user data
 - Subjected to privacy regulations such as GDPR

Data Sources

□ Internal Databases:

- Generated by various services and enterprise applications.
- E. g. Inventory, Customers, Assets
- Need to be safeguarded properly

Types of Data

- First Party Data: Data collected by a company from their users
- Second Party Data: Data collected by a second company from their customers and sold to first company
- Third Party Data: Data collected from the public domain

Data Formats

- Row-Major Format: Consecutive elements in a row are stored next to each other in memory.
 - E. g. CSV format
 - Useful for Transactional Processing
 - Allow faster data writes
- Column-Major Format: Consecutive elements in a column are stored next to each other in memory.
 - E. g. Parquet format
 - Useful for Analytics Processing (e. g. finding mean value of a field)
 - Allow faster reading of a column
 - Pandas DF is column-major
 - Numpy ndarray is row-major by default but order can be specified

Row-major order

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

Column-major order

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

Data Serialization

The process of converting a data structure or object state into a format that can be stored or transmitted and reconstructed later.

Points to consider for deciding serialization format

- Human readability
- Access patterns
- Size of data

Data Serialization

The process of converting a data structure or object state into a format that can be stored or transmitted and reconstructed later.

Major serialization formats:

Format	Binary/Text	Human-readable	Example use cases
JSON	Text	Yes	Everywhere
CSV	<i>Text</i>	Yes	<i>Everywhere</i>
Parquet	<i>Binary</i>	No	<i>Hadoop, Amazon Redshift</i>
Avro	Binary primary	No	Hadoop
Protobuf	Binary primary	No	Google, TensorFlow (TFRecord)
Pickle	Binary	No	Python, PyTorch serialization

Image source: Designing Machine Learning Systems – Chip Guyan

Data Models

Data models describe how data is represented.

Some important types of Data Models are

- Relational Model
- NoSQL Model
 - *Document Model*
 - *Graph Model*

Data Models

❑ Relational Model

- One of the oldest and persistent models
- Invented by Edgar F. Codd in 1970
- Data is organized into relations
- Each relation is a set of tuples
- A relation can be represented as a table

The diagram illustrates a relational table structure. It consists of a grid of horizontal rows and vertical columns. The columns are labeled "Column 1", "Column 2", "Column 3", and "...". A red box highlights the first row, with the text "Tuple (row): unordered" written above it. A blue box highlights the second column, with the text "Column: unordered" written above it.

Column 1	Column 2	Column 3	...

Data Models

❑ Relational Model

- If any attribute is changed multiple rows needs to be updated
- e.g. Change of price
- Useful to convert this into a normalized form

Title	Author	Format	Publisher	Country	Price
Harry Potter	J.K. Rowling	Paperback	Banana Press	UK	\$20
Harry Potter	J.K. Rowling	E-book	Banana Press	UK	\$10
Sherlock Holmes	Conan Doyle	Paperback	Guava Press	US	\$30
The Hobbit	J.R.R. Tolkien	Paperback	Banana Press	UK	\$30
Sherlock Holmes	Conan Doyle	Paperback	Guava Press	US	\$15

Data Models

❑ Relational Model – Normalized Form

Title	Author	Format	Publisher ID	Price
Harry Potter	J.K. Rowling	Paperback	1	\$20
Harry Potter	J.K. Rowling	E-book	1	\$10
Sherlock Holmes	Conan Doyle	Paperback	2	\$30
The Hobbit	J.R.R. Tolkien	Paperback	1	\$30
Sherlock Holmes	Conan Doyle	Paperback	2	\$15

Publisher ID	Publisher	Country
1	Banana Press	UK
2	Guava Press	US

Data Models

❑ Relational Model – Normalized Form

Title	Author	Format	Publisher ID	Price
Harry Potter	J.K. Rowling	Paperback	1	\$20
Harry Potter	J.K. Rowling	E-book	1	\$10
Sherlock Holmes	Conan Doyle	Paperback	2	\$30
The Hobbit	J.R.R. Tolkien	Paperback	1	\$30
Sherlock Holmes	Conan Doyle	Paperback	2	\$15

Publisher ID	Publisher	Country
1	Banana Press	UK
2	Guava Press	US

- In a normalized form data is stored across multiple tables
- Joining multiple tables can be computationally expensive

Data Models

- ❑ Relational Model – Query Language
 - SQL – Most popular query language
 - Declarative language

Data Models

❑ Relational Model – Query Language

- SQL – Most popular query language
- Declarative language:
 - you specify the outputs you want, and the computer figures out the steps needed to get you the queried outputs.
- Imperative language:
 - you specify the steps needed for an action and the computer executes these steps to return the outputs

Data Models

❑ Question: Name some popular relational databases?

Data Models

❑ Answer: Some popular relational databases

- MySQL
- PostgreSQL
- Oracle
- IBM DB2
- Microsoft SQL Server

Data Models

- ❑ NoSQL (Not Only SQL) Data Model
 - Relational Data Model demands that data follows a strict schema
 - Schema management can be painful
 - Two main NoSQL Data Models:
 - Document Model
 - Graph Model

Data Models

❑ Document Type NoSQL Data Model

- Built around the concept of a "Document"
- A "Document" is a single continuous string
- Encoded as JSON, XML or BSON
- All documents in a database are assumed to be encoded by the same format
- Each document is represented by a unique key
- Analogy with relational tables:
 - A single document is similar to a row
 - Collection of documents is similar to a table

Data Models

□ Example of Document Type NoSQL Data Model

Example 3-1. Document 1: harry_potter.json

```
{  
    "Title": "Harry Potter",  
    "Author": "J .K. Rowling",  
    "Publisher": "Banana Press",  
    "Country": "UK",  
    "Sold as": [  
        {"Format": "Paperback", "Price": "$20"},  
        {"Format": "E-book", "Price": "$10"}  
    ]  
}
```

Example 3-2. Document 2: sherlock Holmes.json

```
{  
    "Title": "Sherlock Holmes",  
    "Author": "Conan Doyle",  
    "Publisher": "Guava Press",  
    "Country": "US",  
    "Sold as": [  
        {"Format": "Paperback", "Price": "$30"},  
        {"Format": "E-book", "Price": "$15"}  
    ]  
}
```

Example 3-3. Document 3: the_hobbit.json

```
{  
    "Title": "The Hobbit",  
    "Author": "J.R.R. Tolkien",  
    "Publisher": "Banana Press",  
    "Country": "UK",  
    "Sold as": [  
        {"Format": "Paperback", "Price": "$30"},  
    ]  
}
```

Data Models

- ❑ Document Type NoSQL Data Model – Advantages
 - Schema-less
 - Faster creation and read
 - Better locality

- ❑ Document Type NoSQL Data Model – Disadvantages
 - Difficult to make analytics computations
 - Consistency limitations
 - Atomicity weakness

Data Models

❑ Question: Name some popular NoSQL document databases?

Data Models

□ Answer: Some popular NoSQL document databases

- MongoDB
- Amazon DocumentDB
- CouchDB

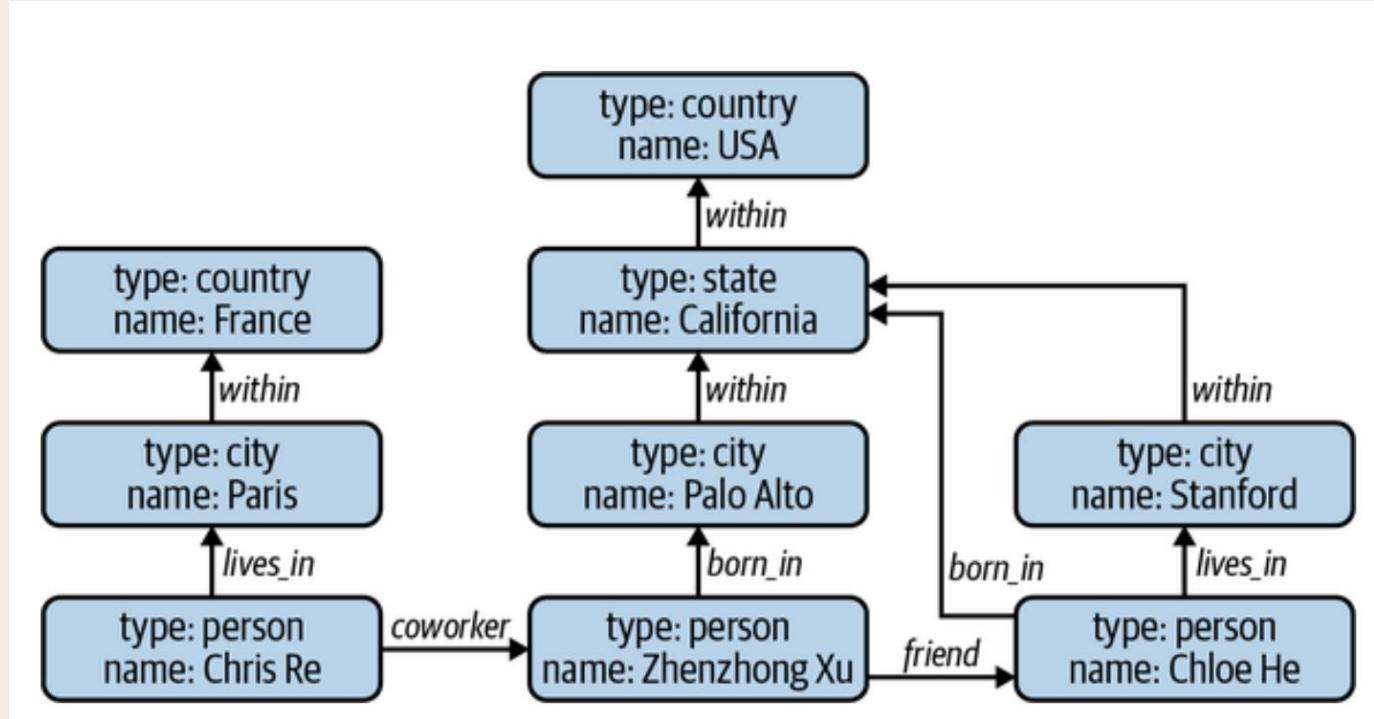
Data Models

❑ Graph Type NoSQL Data Model

- Built around the concept of a "Graph"
- A "Graph" consists of Nodes and Edges
- Edges represent relationship between nodes
- Importance:
 - Document Data Model: content of each data item
 - Graph Data Model: relationship between each data item

Data Models

- Graph Type NoSQL Data Model – Example



- Graph databases are queried using Graph Query Languages:
- Homework exercise: read further information on GQL from https://en.wikipedia.org/wiki/Graph_Query_Language

Data Models

❑ Question: Name some popular Graph databases?

Data Models

□ Answer: Some popular Graph databases

- Neo4J
- OrientDB
- Amazon Neptune

Data Processing Types

Databases are optimized for two types of workloads

- Transactional Processing
- Analytical Processing

Transactional Processing

- ❑ Transaction refers to any kind of action:
 - Ordering a Pizza
 - Uploading a YouTube video
 - Booking an Uber Cab
- ❑ Transactions are inserted into a DB as they are generated
- ❑ Occasionally transactions are updated if some fields are changed
- ❑ Also known as Online Transaction Processing or OLTP

Transactional Processing

- ❑ Transactional databases are designed to satisfy ACID properties:
 - **Atomicity:**
 - To guarantee that all the steps in a transaction are completed successfully as a group.
 - If any step in the transaction fails, all other steps must fail also
 - e. g. If user payment fails, entire Uber booking transaction is cancelled.
 - **Consistency:**
 - To guarantee that all the transactions coming through must follow predefined rules.
 - e. g. Only authorized users can book an Uber cab

Transactional Processing

- ❑ Transactional databases are designed to satisfy ACID properties:
 - Isolation:
 - To guarantee that two transactions happen at the same time as if they were isolated.
 - e. g. Two passengers cannot book the same Uber cab at the same time.
 - Durability:
 - To guarantee that once a transaction has been committed, it will remain committed even in the case of a system failure.
 - e. g. after you've ordered a ride and even if you lost network connection, you still want your ride to come.

Transactional Processing

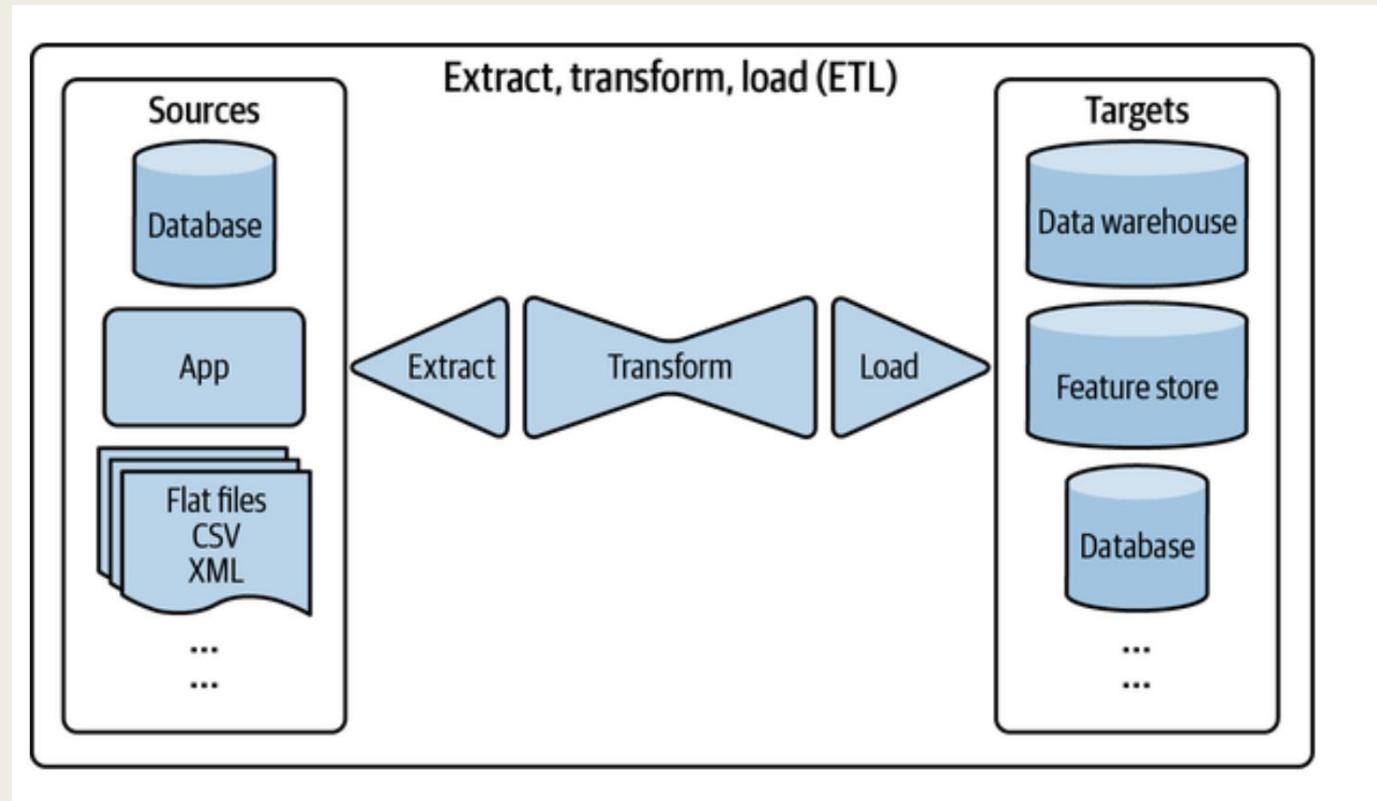
- ❑ ACID Requirements can be too restrictive for many use cases
- ❑ Instead, BASE requirements are used in many situations
- ❑ BASE - Basically Available, Soft state, and Eventual consistency

Analytical Processing

- These databases are optimized for making aggregations of a column across rows
 - e. g. *To find average price of an Uber ride in a month at a particular city*
- Also known as *Online Analytics Processing or OLAP*

Extract, Transform and Load (ETL)

- ❑ Data is first **Extracted** from different sources
- ❑ It is then **Transformed** into the desired Schema
- ❑ Finally **Loaded** to the target database



Breakout Session

- Discuss what type of data processing (Transactional or Analytics) is the major type in your project
- Based on this arrive at what type of database you should be using for your project
- Also what type of data pipeline your project would be required