



CST8507: NATURAL LANGUAGE PROCESSING

WEEK#4
WORD EMBEDDING

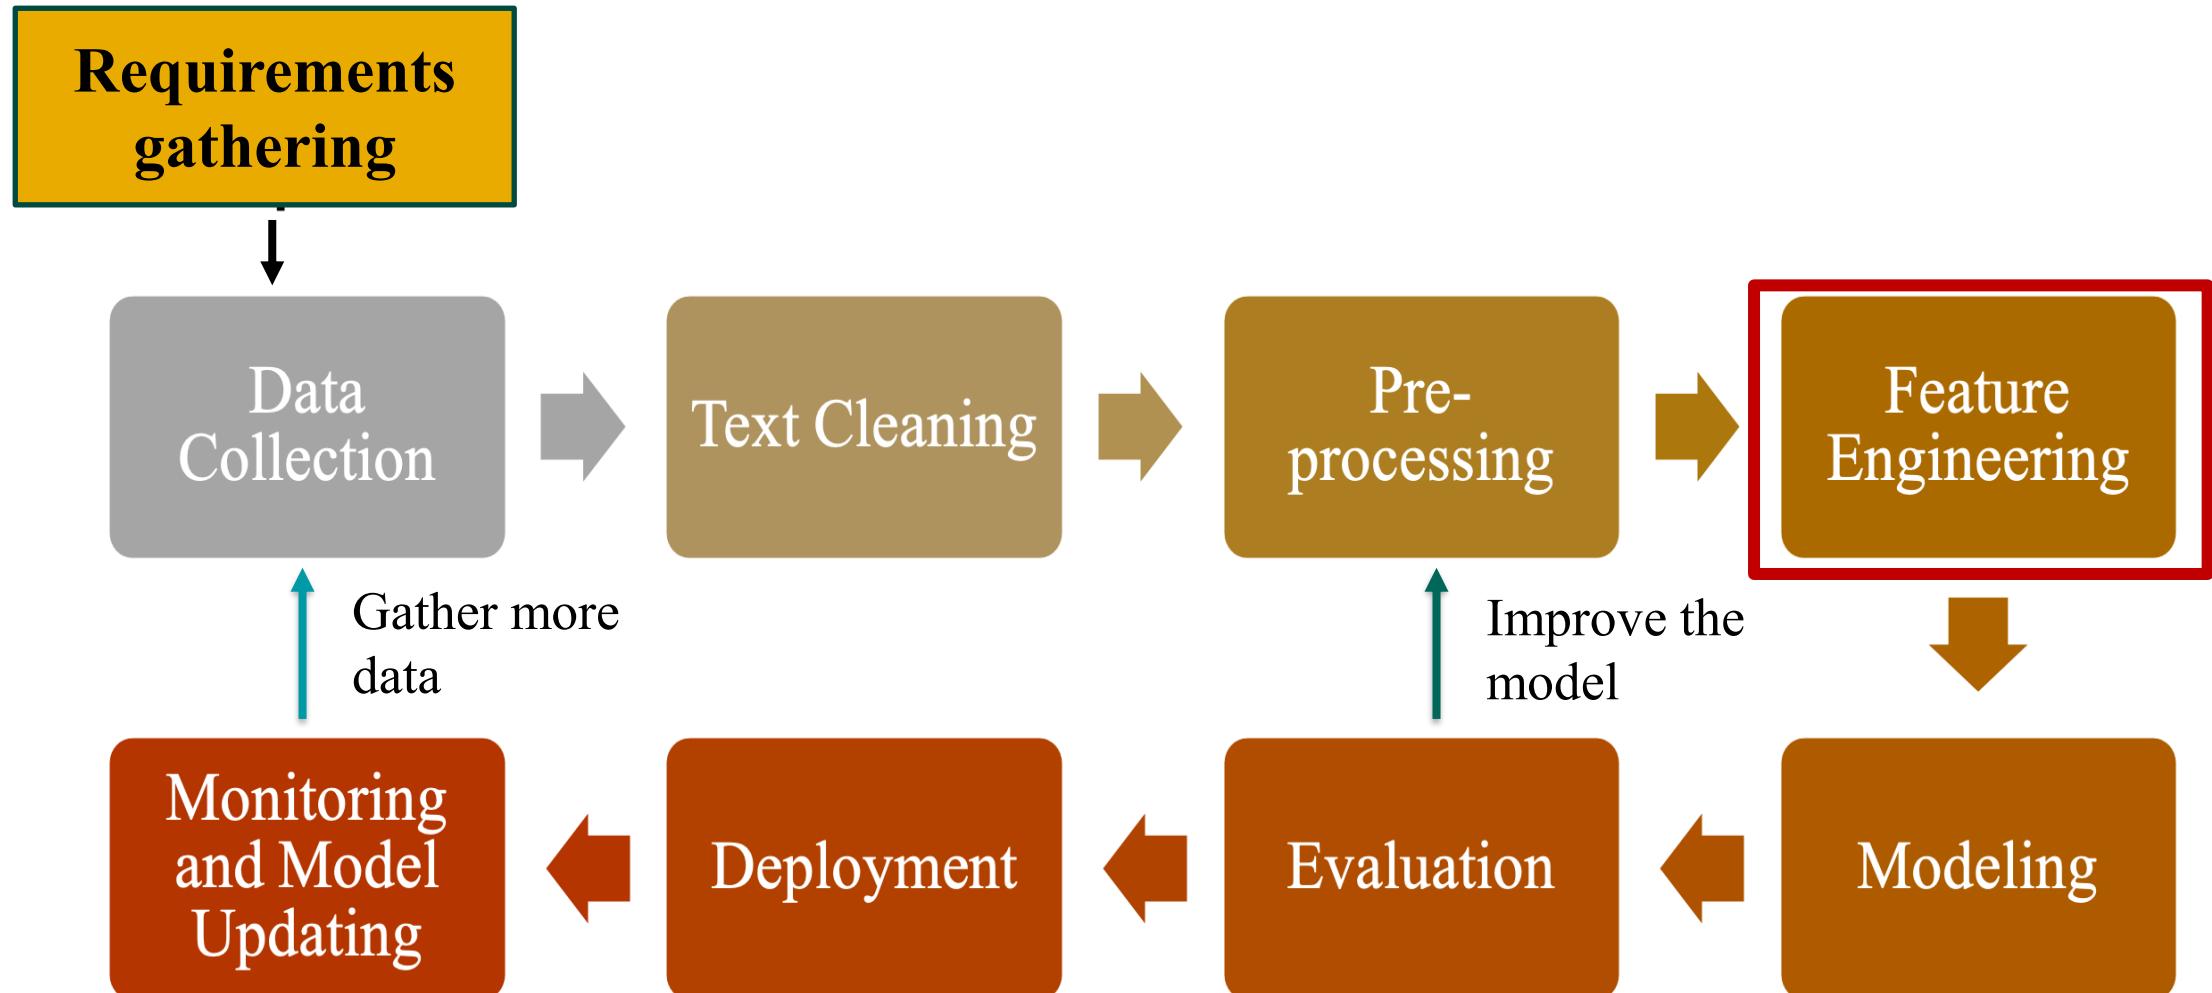
DEVELOPED BY
HALA OWN, PH.D.

Lesson Agenda

- **Prediction based Text representation (word Embedding)**
 - CBOW, Skip-Gram and SGNS
 - Word2Vec
 - FastText
- **Count-Based / Matrix Factorization Methods**
 - Glove



NLP Development Life Cycle



Text Representation Techniques

- **Frequency based Text representation:**
 - One-Hot Encoding
 - Bag of Words
 - Bag of N-Grams
 - TF-IDF
- **Prediction based Text representation (word Embedding)**
- **Universal Text Representations**



Comparing Feature Representations For Audio, Image And Text

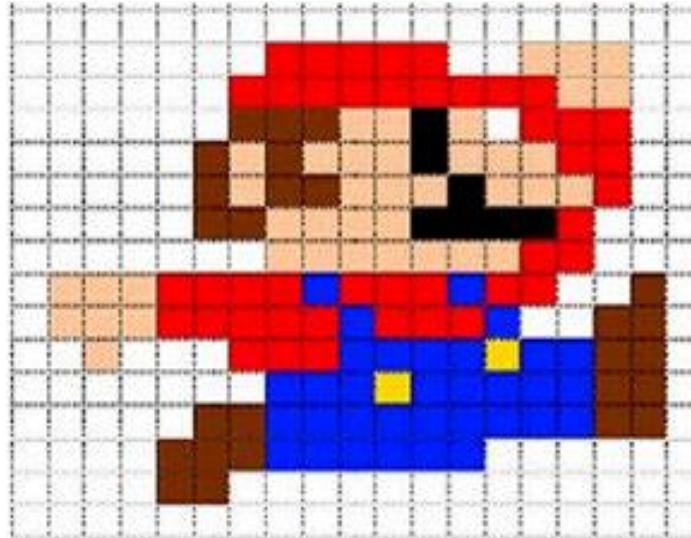
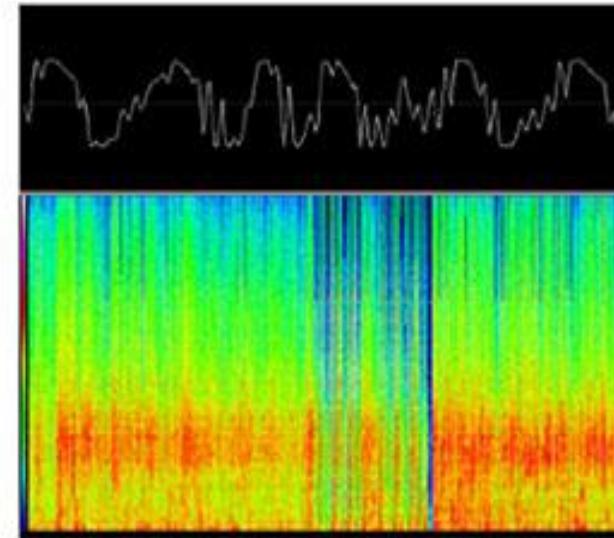


IMAGE PIXELS (DENSE)



AUDIO SPECTROGRAM (DENSE)

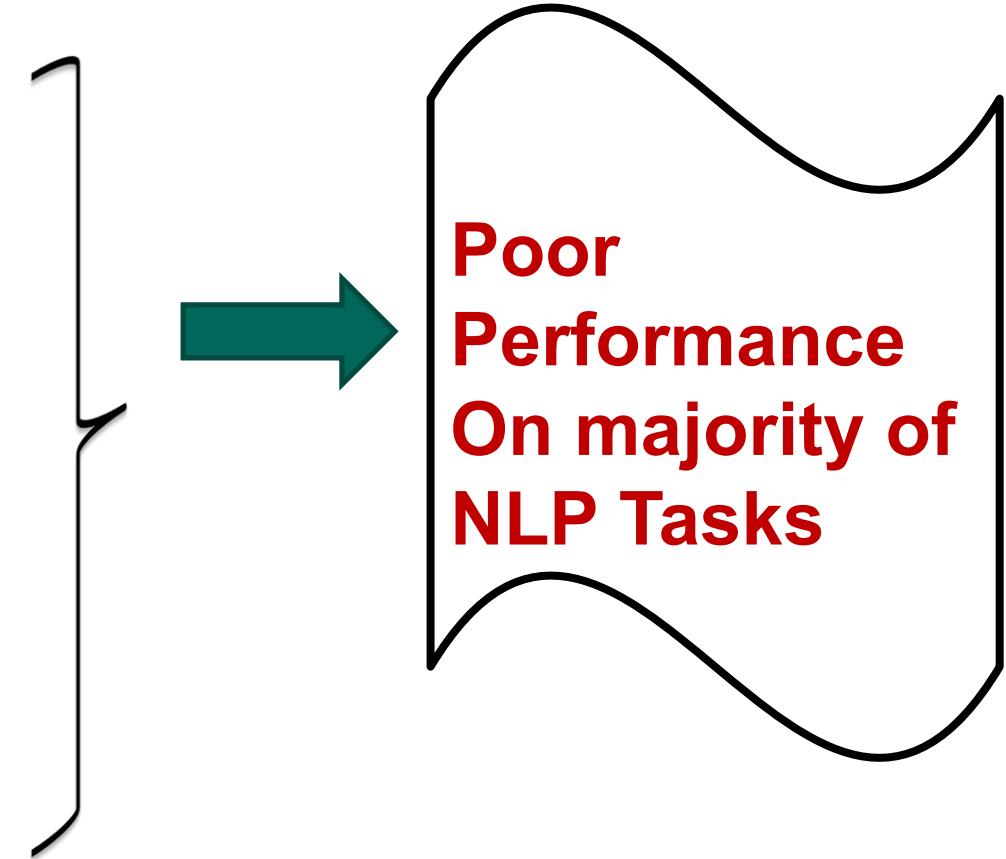


WORD VECTORS (SPARSE)



Frequency based Text representation: Limitation

- High-dimensional representation
- sparse
- OOV words
- Lack of semantic meaning



Relation between Word Senses: Word Similarity

- Cat is not a synonym of dog, but cats and dogs are certainly similar words
- A semantic field is a set of words which cover a particular semantic domain
 - Restaurants: waiter, menu, plate, food, chef
 - Houses: door, roof, kitchen, family, bed

One way of getting values for word similarity is to ask humans to judge how similar one word is to another



WordNet

a large lexical database of English words, where words are grouped into sets of synonyms called **synsets**. These synsets are connected by various semantic relationships, such as synonymy, hypernymy, and hyponymy,etc.

WordNet Search - 3.1
- [WordNet home page](#) - [Glossary](#) - [Help](#)

Word to search for:

Display Options: ▾

Key: "S:" = Show Synset (semantic) relations, "W:" = Show Word (lexical) relations
Display options for sense: (gloss) "an example sentence"

Noun

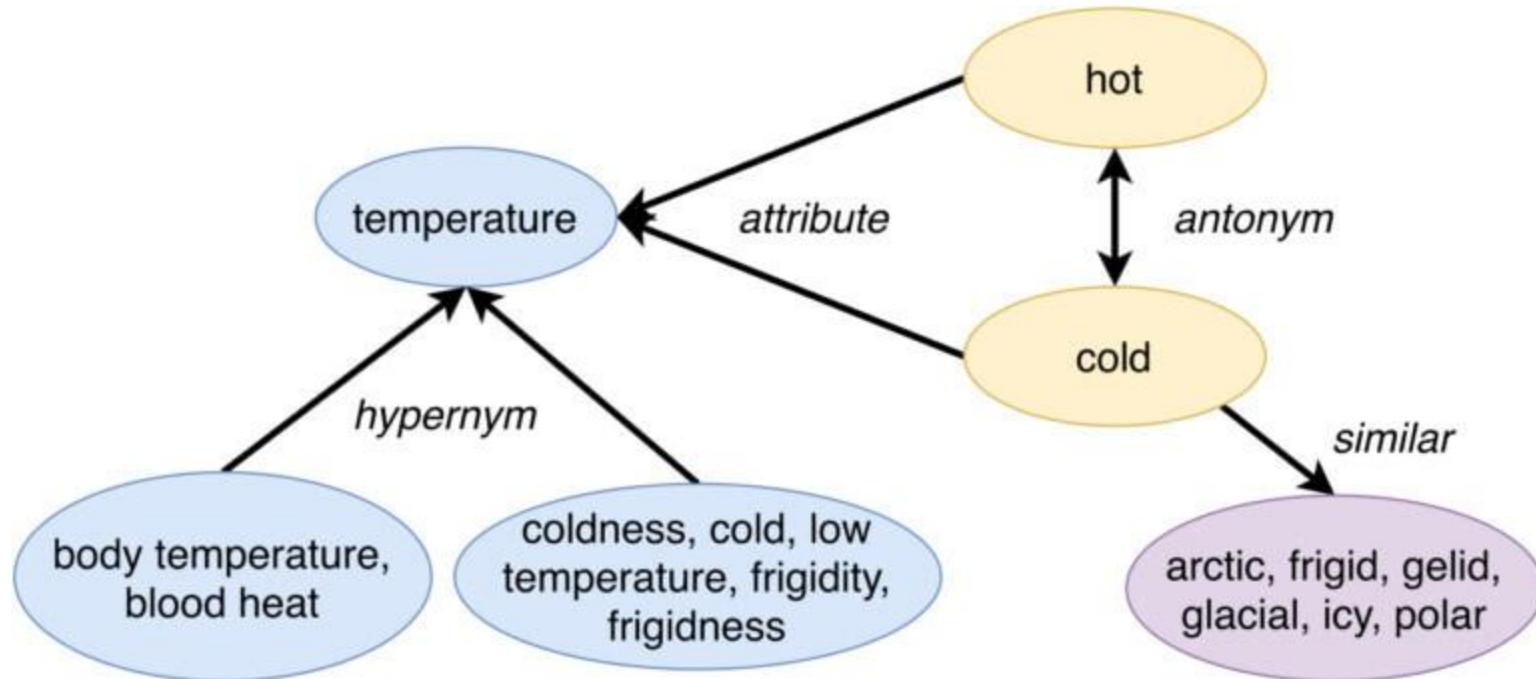
- [S: \(n\) wordnet](#) (any of the machine-readable lexical databases modeled after the Princeton WordNet)
- [S: \(n\) WordNet, Princeton WordNet](#) (a machine-readable lexical database organized by meanings; developed at Princeton University)

<https://wordnet.princeton.edu/>



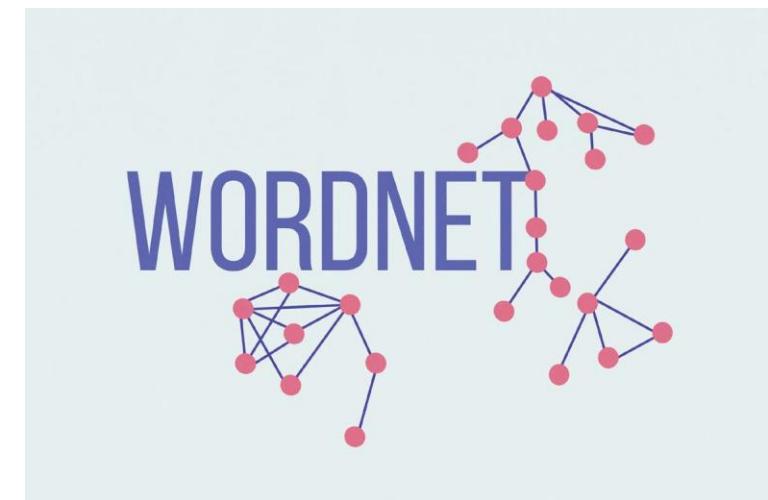
WordNet...

Database of lexical relations for English

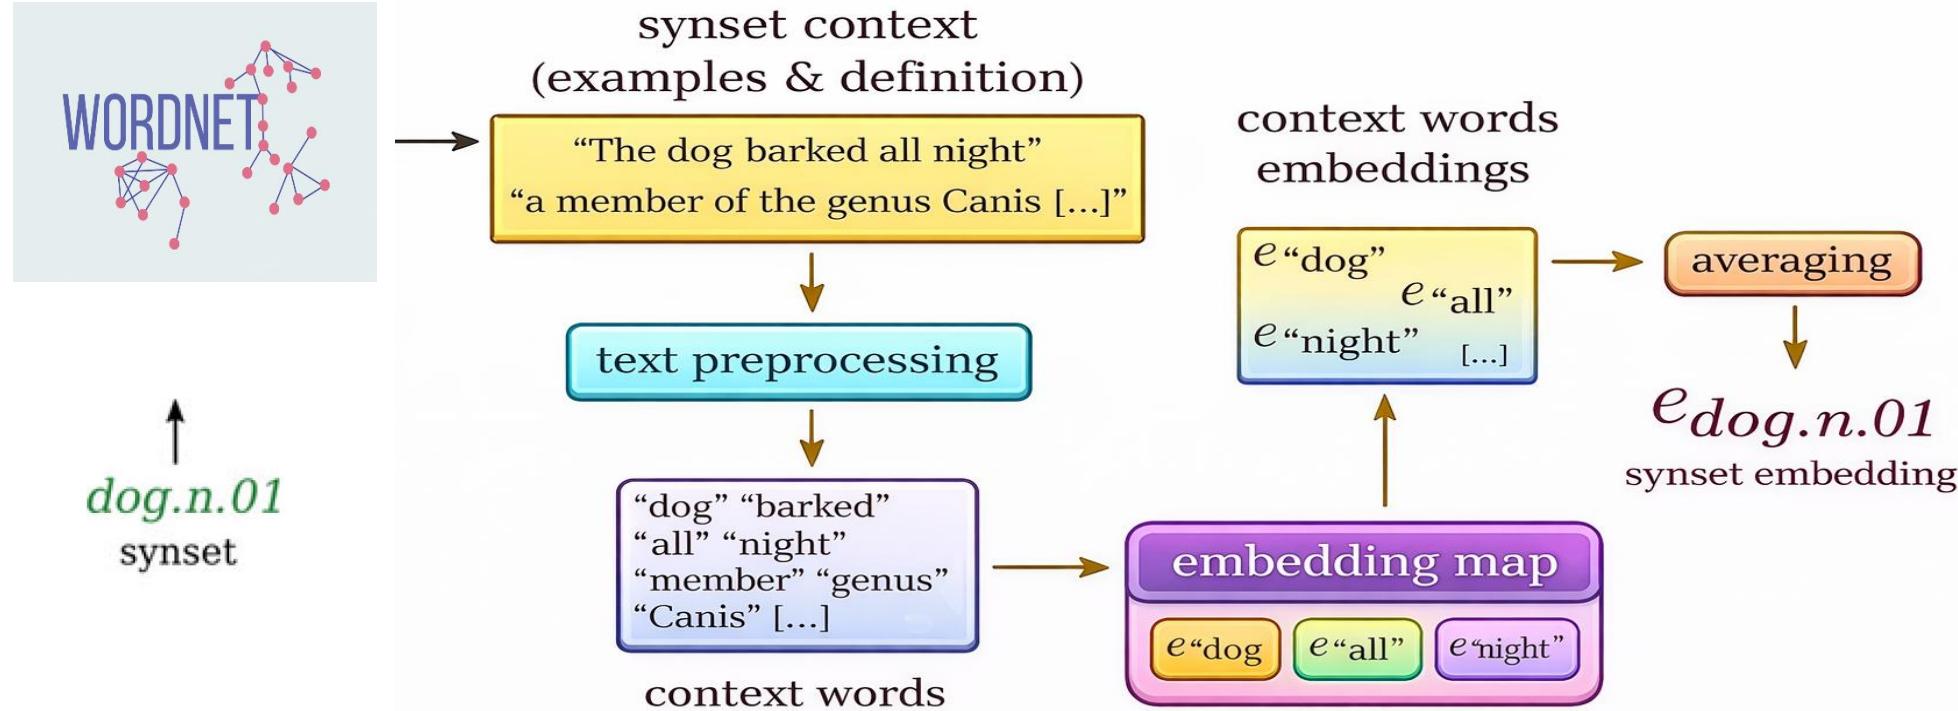


WordNet – Important Concepts

- ❖ **Synset:** A set of synonyms that share a common meaning.
- ❖ **Hypernym:** A general term that encompasses more specific terms (e.g., "animal" is a hypernym of "dog").
- ❖ **Hyponym:** A specific term within a broader category (e.g., "dog" is a hyponym of "animal").
- ❖ **Meronym:** A term that denotes a part of something (e.g., "wheel" is a meronym of "car").
- ❖ **Holonym:** A term that denotes a whole of which the meronym is a part (e.g., "car" is a holonym of "wheel").
- ❖ **Antonym:** Words that have opposite meanings (e.g., "hot" and "cold").
- ❖ **Troponym:** A verb that denotes a specific manner of doing something (e.g., "run" is a troponym of "move").
- ❖ **Entailment:** A relationship where one verb implies another (e.g., "snore" entails "sleep").



WordNet Applications in NLP Tasks: Semantic Test presentation



The context of each synset is tokenized into words, with each word mapped to a vector representation via the learned embedding matrix. The synset vector is the centroid produced by averaging all context word embeddings.

WordNet Applications in NLP Tasks: Query Expansion

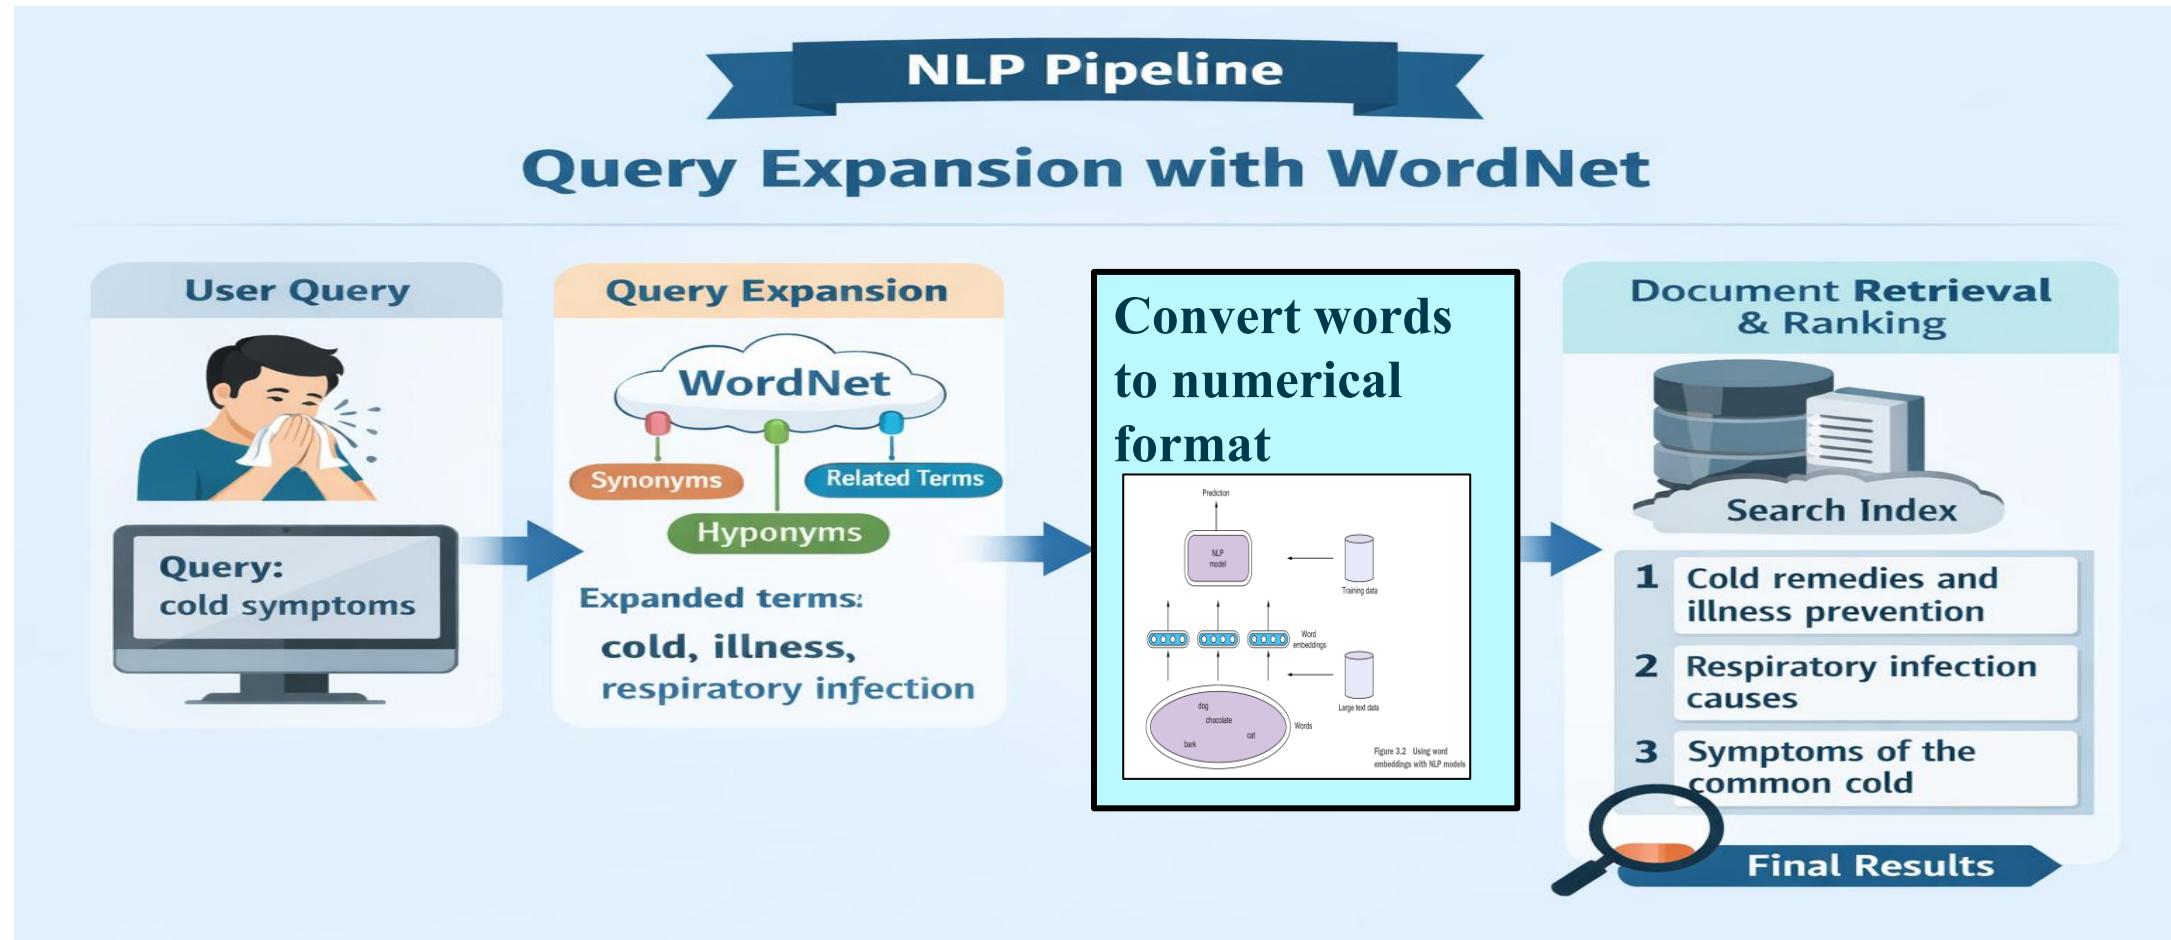


Image generated by ChatGPT with some updates

WordNet: Limitations

- ❑ Limited Coverage and Static Nature
- ❑ Not Computational
- ❑ Domain Specificity
- ❑ Language Limitation
- ❑ Manual Curation Challenges



Key Terms

- ***Distributional similarity***: the meaning of a word can be understood from the context
- ***Distributional hypothesis***: words that occur in similar contexts have similar meanings.



Vector Semantics(Word Embedding)

Computational model that learn the **linguistic units** (words, phrases, or documents) representations based on distributional properties of these units in a large corpus.

- Representation **linguistic units** as **vectors in a multi-dimensional space**.
- Encoding semantic information using mathematical vectors.
- Standard way to represent word meaning in¹⁵ NLP.



Word Embedding

Representation of words as vectors of numbers in a high-dimensional space. It captures **semantic** and **contextual information** about the word.

Input: large number of corpus, Vocabulary V , and vector of dimension d

output

$$f: V \rightarrow R^d$$

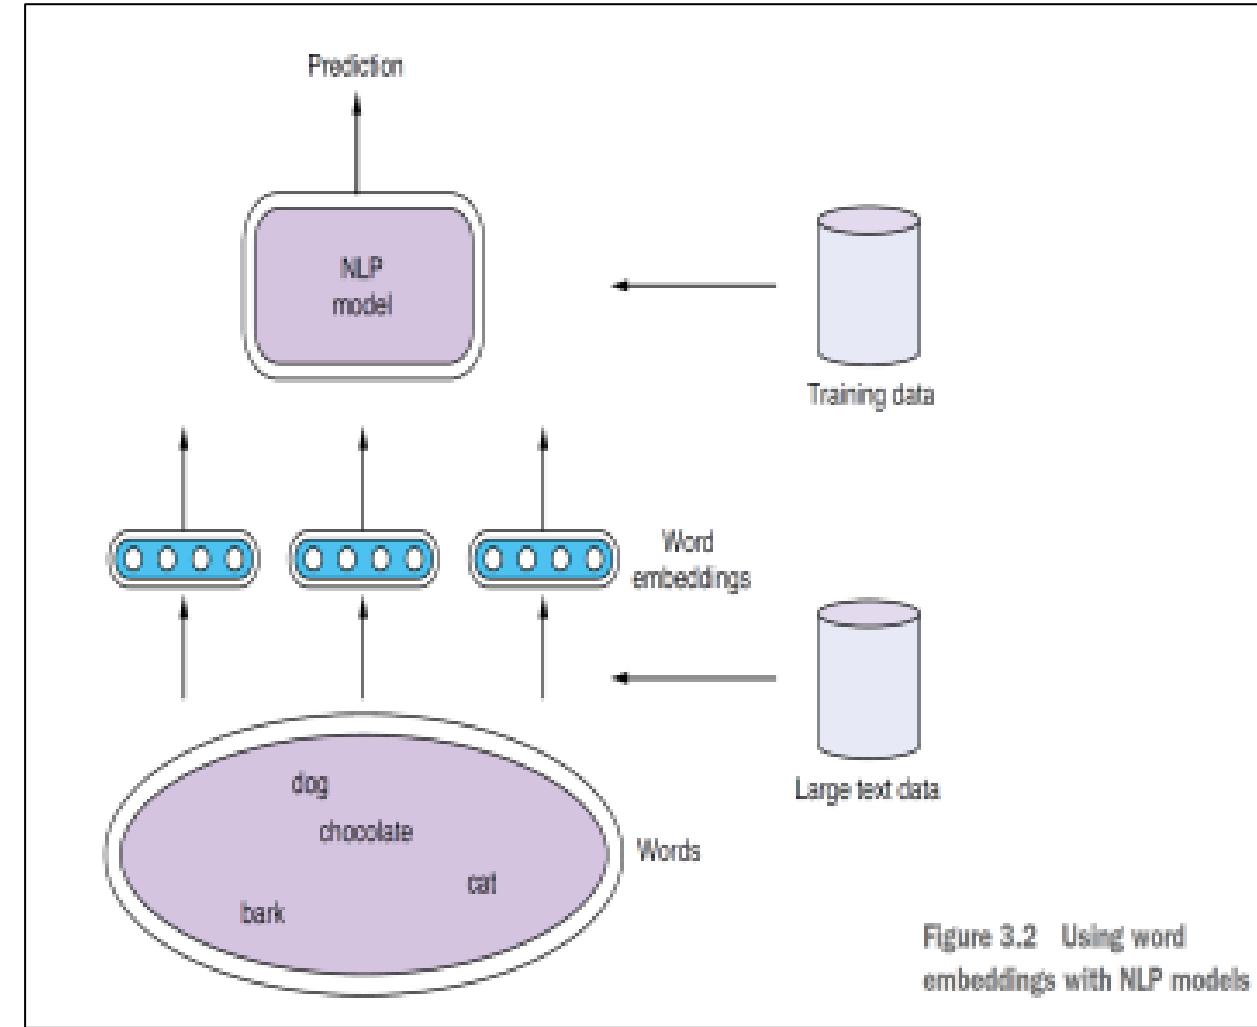


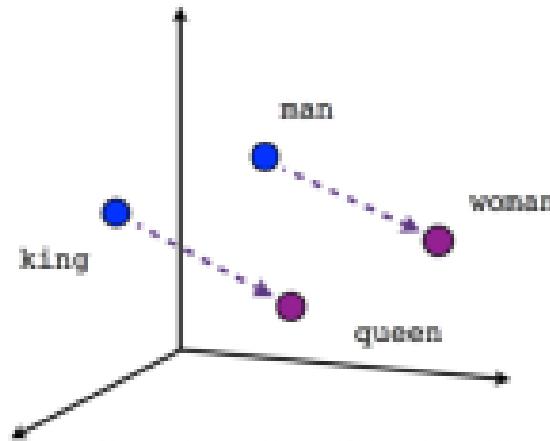
Figure 3.2 Using word embeddings with NLP models



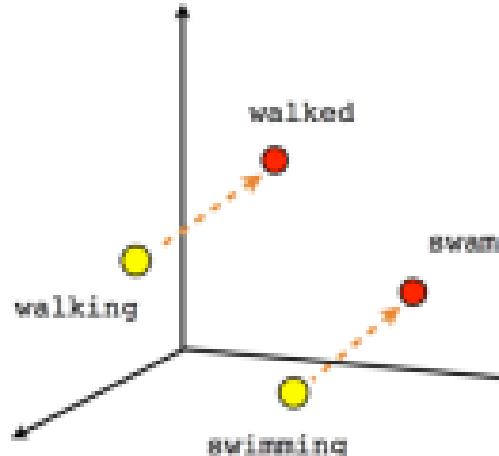
Image source: real world natural language processing book

ALGONQUIN
COLLEGE

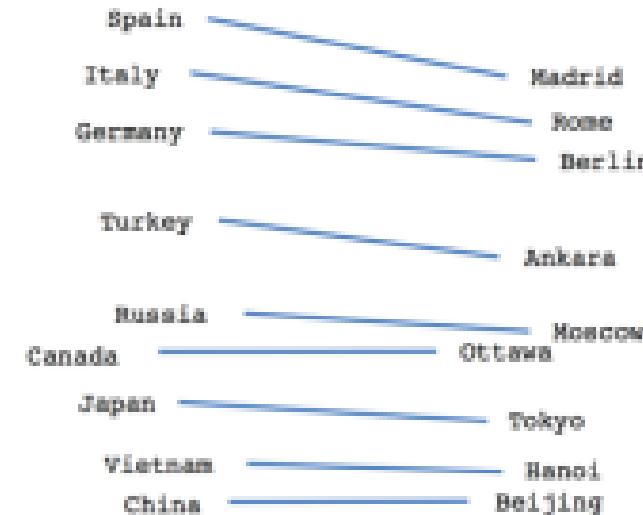
Word Embedding... Analogy



Male-Female



Verb tense



Country-Capital

`vector('king') - vector('man') + vector('woman') ≈ vector('queen')`

Prediction based Text representation

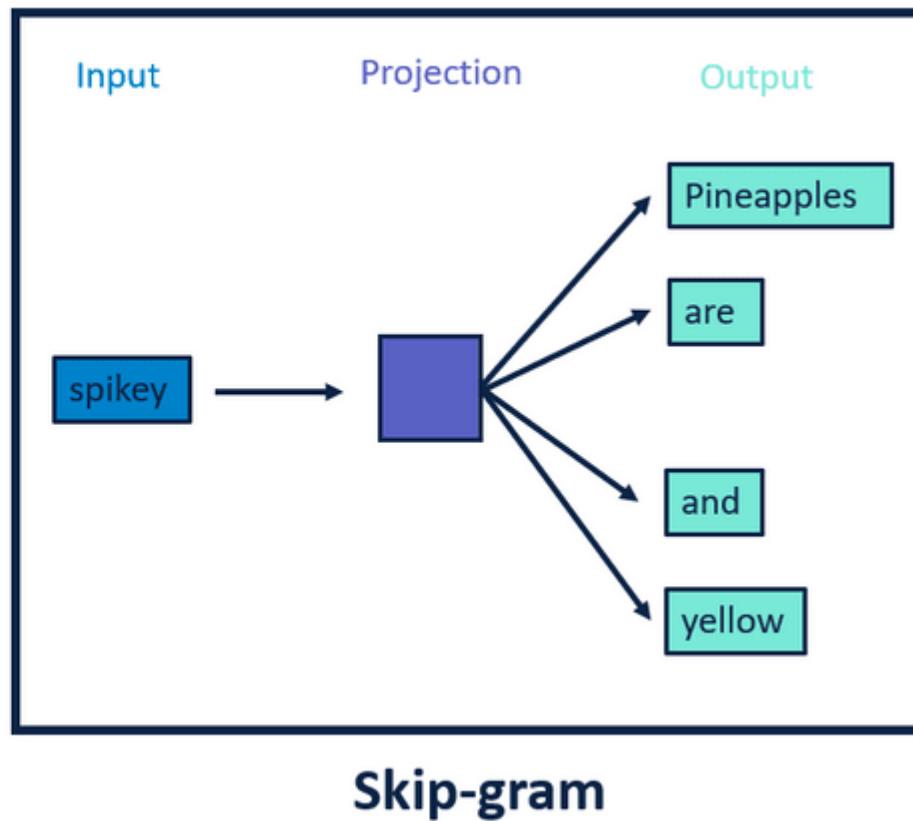
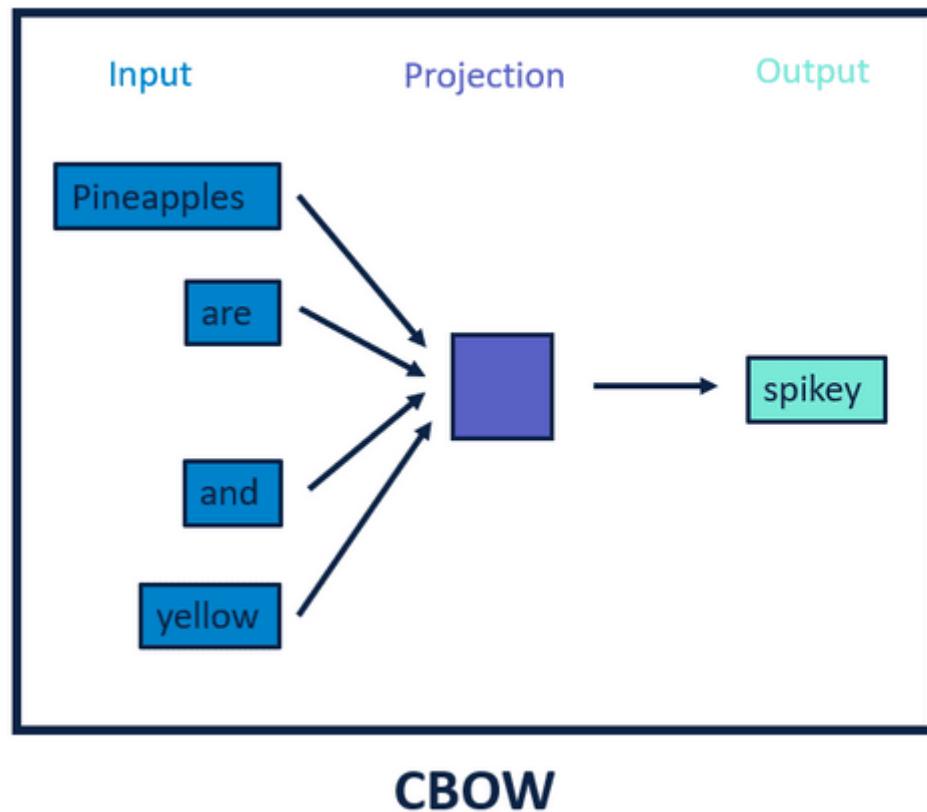
Big idea: self-supervision, Bengio et al. (2003) and Collobert et al. (2011)

- Popular embedding method
- Very **fast to train**
- Code available on the web
- **Predict** rather than **count**



Prediction Based Text Representation(Word2Vec)

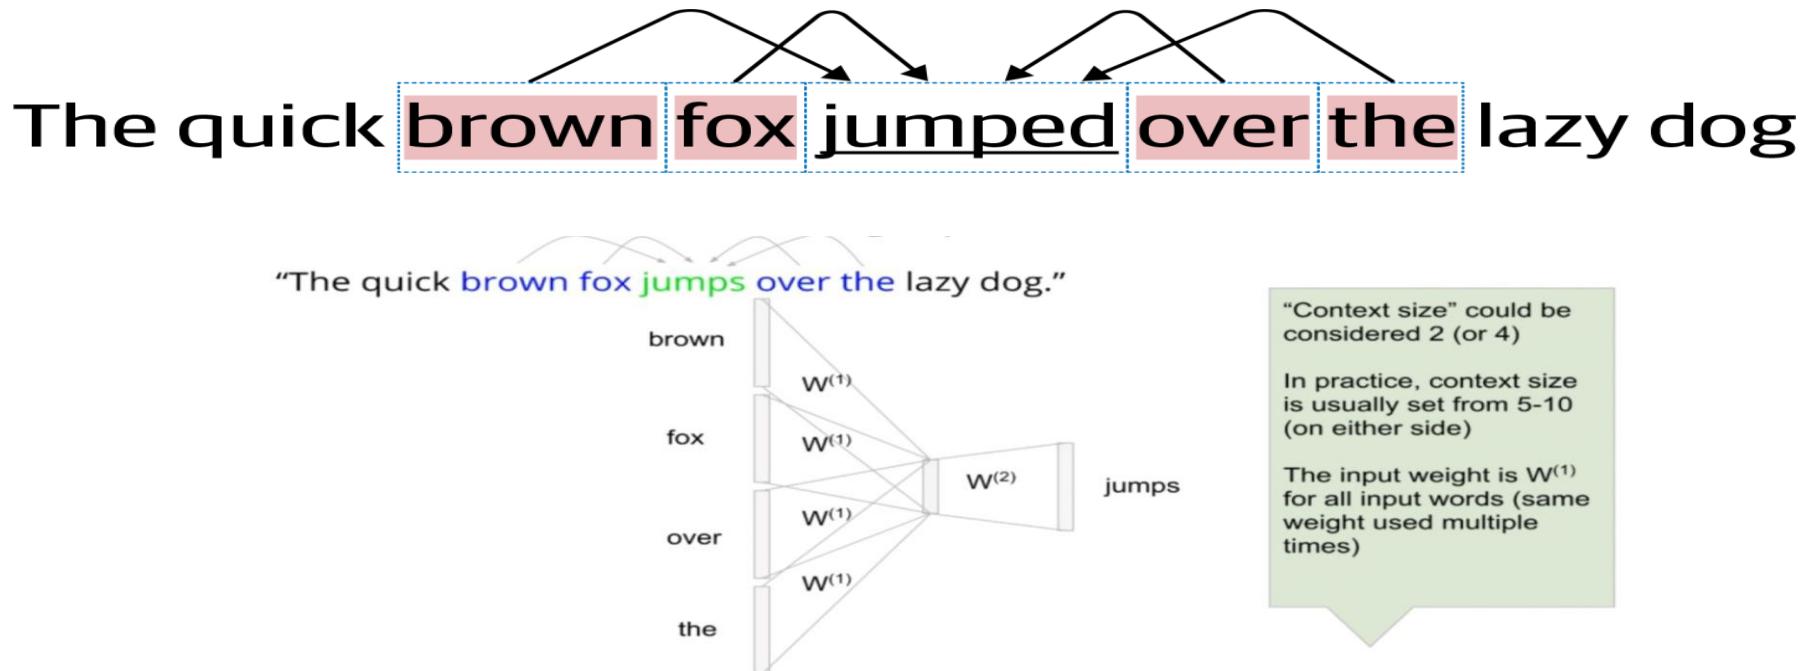
Word Embedding (creating dense vector representations of words)



Source: <https://dataaspirant.com/word-embedding-techniques-nlp/>

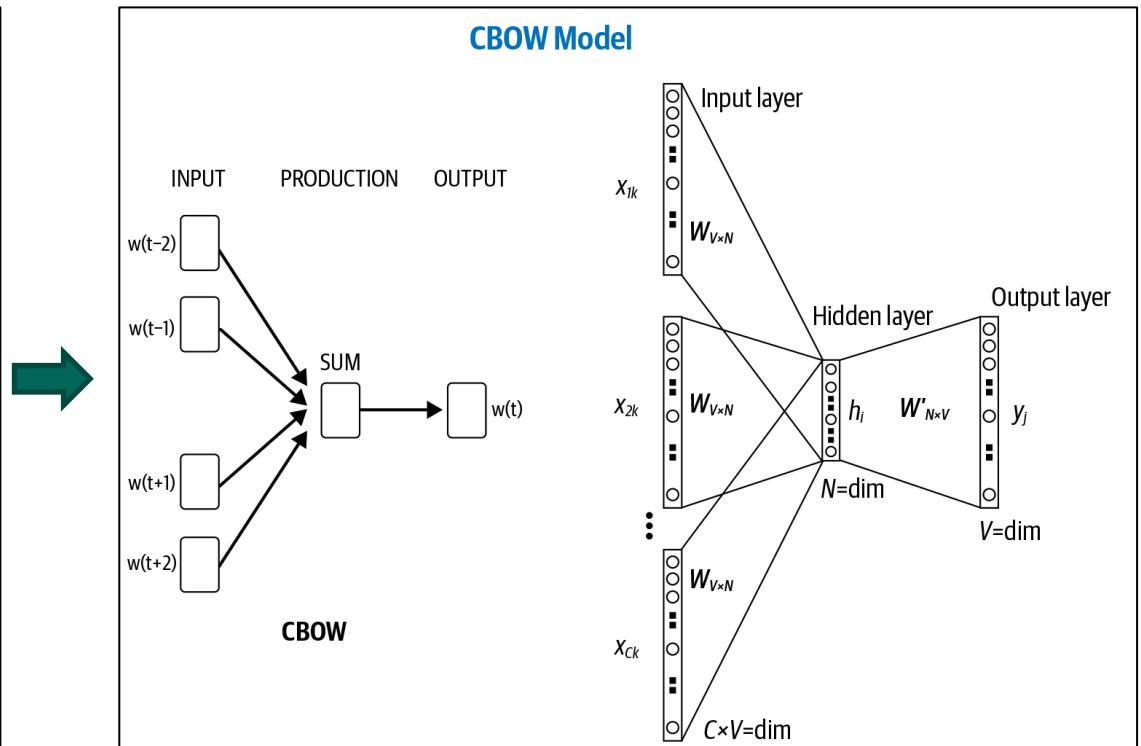
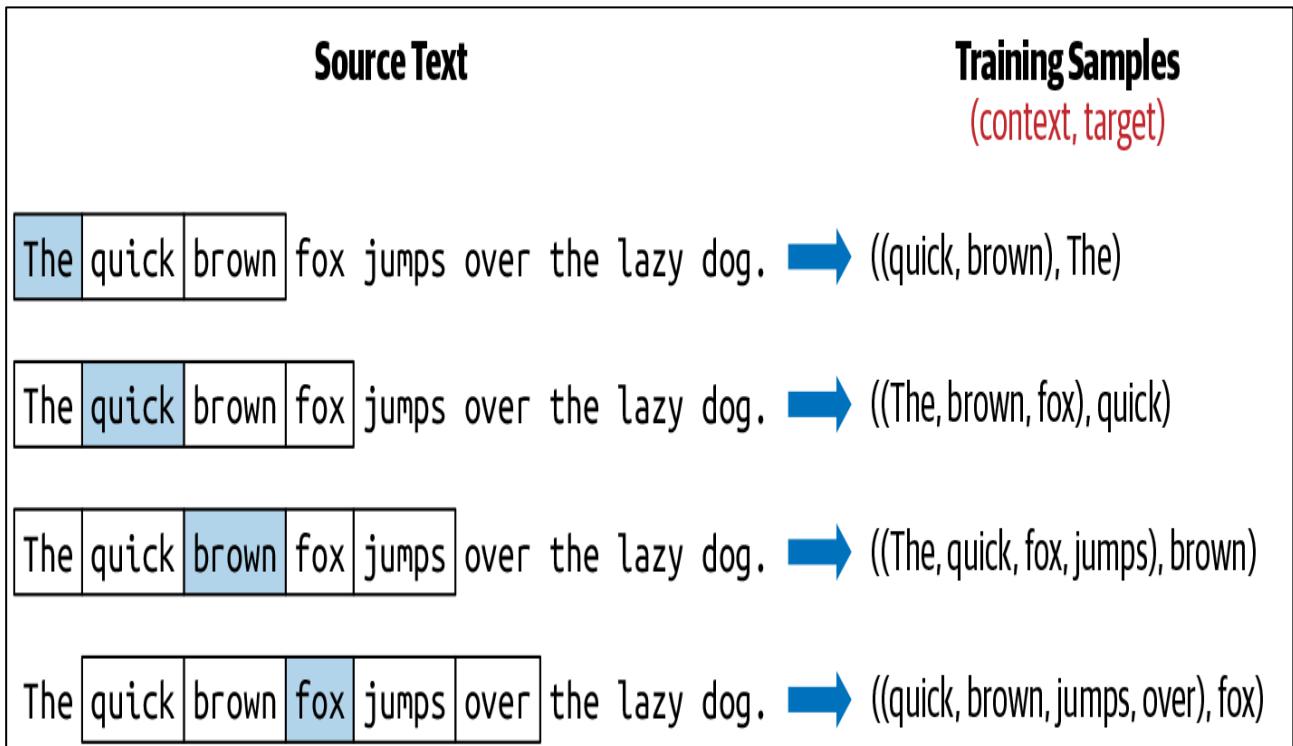
CBOW

- Goal: Predict the **middle** word given the **words of the context**



CBOW..

Window size=2

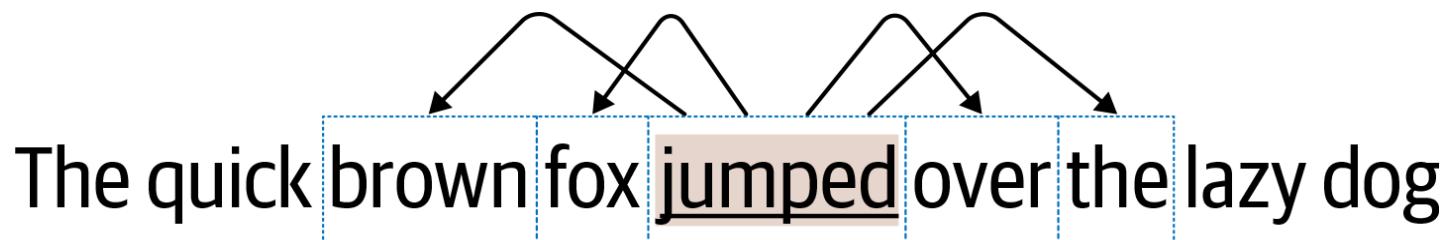


CBOW: Simple Example

		Context										Input-Hidden Weight				Hidden Activation			
												1	2	3	4	5	6	7	8
C1	this	0	1	0	0	0	0	0	0	0	0	13	14	15	16	17	18	19	20
C2	corpus	0	0	0	0	1	0	0	0	0	0	17	18	19	20	33	34	35	36
C3	context	0	0	0	0	0	0	0	0	1	0	21	22	23	24	Average hidden Activation			
												25	26	27	28	29	30	31	32
												33	34	35	36	18.333333333 19.333333333 20.333333333 21.333333333			
												37	38	39	40				

Skip-gram

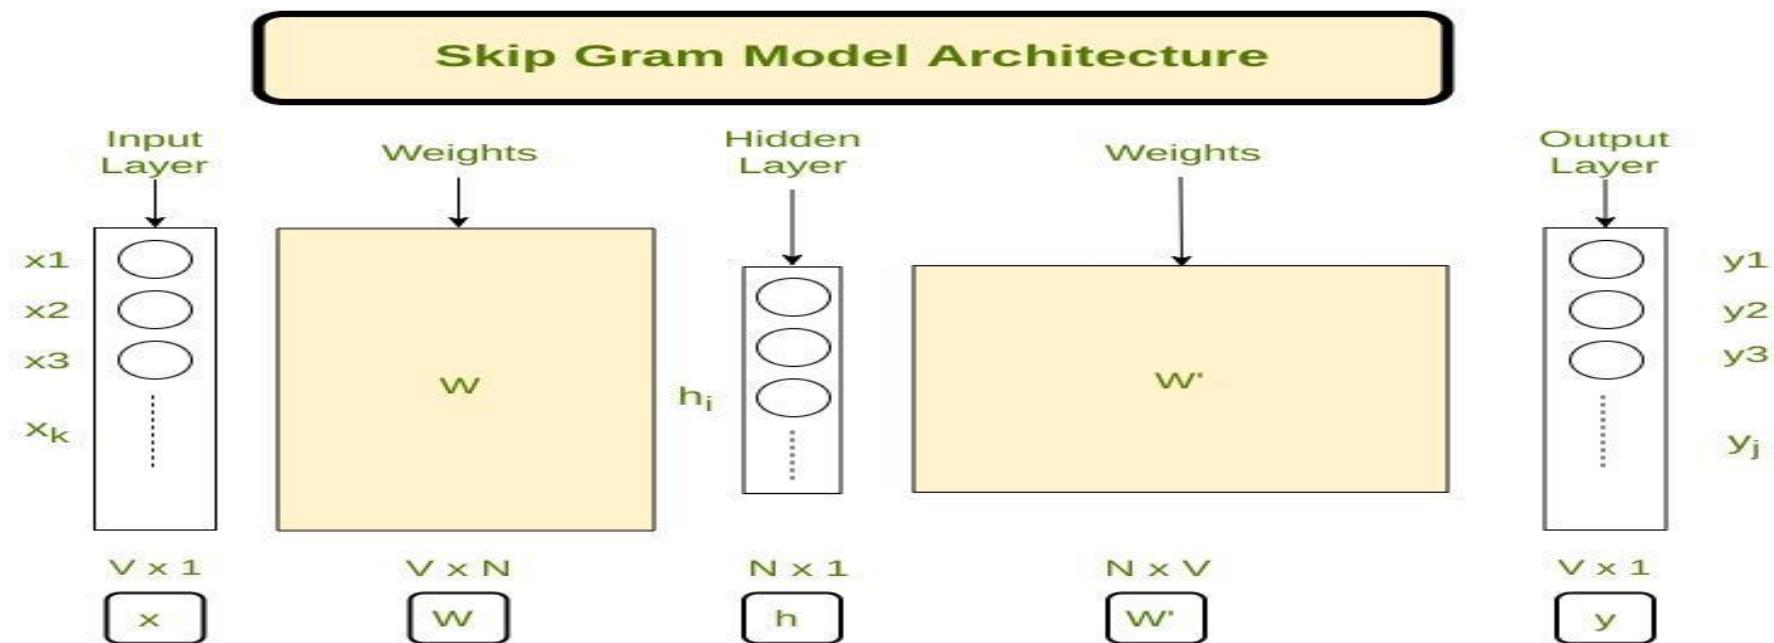
- Goal: Predict the **context** words given the **middle word**



Source Text	Training Samples
The quick brown fox jumps over the lazy dog. →	(the, quick) (the, brown)
The quick brown fox jumps over the lazy dog. →	(quick, the) (quick, brown) (quick, fox)
The quick brown fox jumps over the lazy dog. →	(brown, the) (brown, quick) (brown, fox) (brown, jumps)
The quick brown fox jumps over the lazy dog. →	(fox, quick) (fox, brown) (fox, jumps) (fox, over)

Skip-gram...

the training objective is to minimize the summed prediction error across all context words in the output layer.



Skip-gram: Example

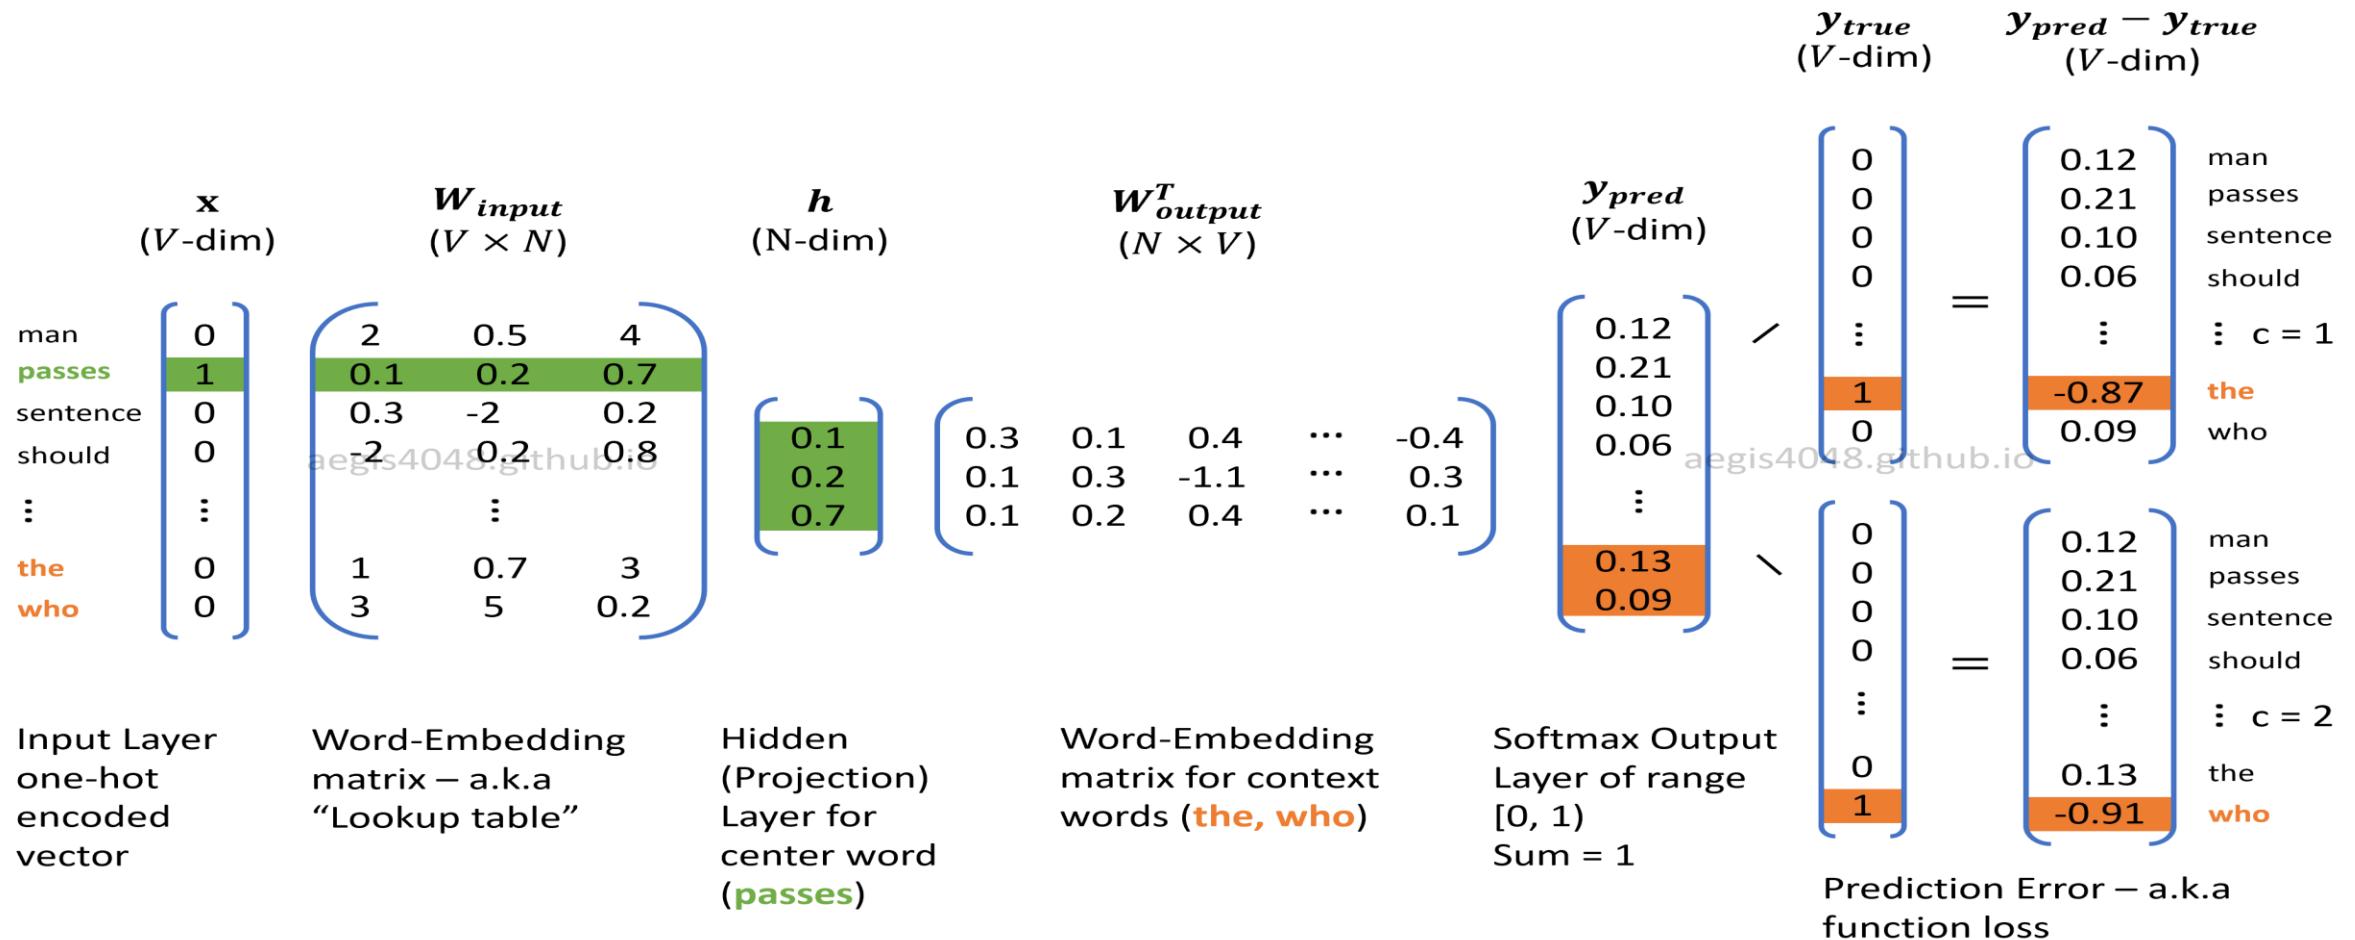


Image source:
https://aegis4048.github.io/demystifying_neural_network_in_skip_gram_language_modeling

Skip-gram Prediction: Example

the cat sat on the mat



context size = 2



Skip-gram Prediction :Example...



context size = 2



Skip-gram Prediction: :Example...

the cat sat on the mat



context size = 2



ALGONQUIN
COLLEGE

Skip-gram Prediction :Example...

the cat sat on the mat



context size = 2



Skip-gram Prediction :Example...

the cat sat on the mat



context size = 2



ALGONQUIN
COLLEGE

Skip-gram Prediction :Example...

the cat sat **on the mat**



context size = 2



Skip-gram vs CBOW

- CBOW is comparatively **faster** to train than skip-gram and **better for frequently occurring words**
- Skip-gram is slower but works well for **smaller amount of data**
- CBOW is an **easier** classification problem than Skip-gram



Skip-gram Negative Sampling(SGNS): Approach

1. Treat the target word t and a neighboring context word c as **positive examples**.
2. Randomly sample other words in the lexicon to get negative examples
3. Use **logistic regression** to train a classifier to distinguish those two cases
4. Use the learned weights as the embeddings



SGNS : how to learn vectors

- Given the set of positive and negative training instances, and an initial set of embedding vectors
- The goal of learning is to adjust those word vectors such that we:
 - **Maximize** the similarity of the **target word, context word** pairs (w, c_{pos}) drawn from the positive data
 - **Minimize** the similarity of the (w, c_{neg}) pairs drawn from the negative data.



SGNS : how to learn vectors...

- Training sentence:

... lemon, a tablespoon of **apricot** jam a pinch

...

positive examples +	
t	c
apricot	tablespoon
apricot	of
apricot	preserves
apricot	or

c1

c2

t

c3

c4

negative examples -

t	c	t	c
apricot	aardvark	apricot	twelve
apricot	puddle	apricot	hello
apricot	where	apricot	dear
apricot	coaxial	apricot	forever



Pretrained Word Embeddings Models

- **Word2vec** (Mikolov et al.) 2013
- <https://code.google.com/archive/p/word2vec/>
- **Fasttext** <http://www.fasttext.cc/> 2016
- **Glove** (Pennington, Socher, Manning) 2014
- <http://nlp.stanford.edu/projects/glove/>



Google's Word2Vec

- **Gensim package** : Google's *pre-trained* Word2Vec model in Python.
- This model is trained on the vocabulary of **3 million** words and phrases from 100 billion words of the Google News dataset.
- The vector length for each word is **50,100,300**.



Google's Word2Vec

- Install genism library
`conda install -c conda-forge genism`
- Genism word2vec Model Training

```
model = Word2Vec(text, min_count=1, vector_size= 50, window =5,  
sg = 1, negative=5)
```



Demo

- inclassCode



The GloVe (Global Vector for word presentation)

- Unsupervised learning model that can be used to obtain dense word vectors.
- invented in Stanford by Pennington et al.
<https://nlp.stanford.edu/projects/glove/>
- <https://github.com/stanfordnlp/GloVe>



Glove Algorithm

Step—1: Collect a large dataset of words and their co-occurrences. This can be a text corpus, such as documents or web pages.

Step—2: Preprocess the dataset by tokenizing the text into individual words and filtering out rare or irrelevant words.

Step—3: Construct a co-occurrence matrix that counts the number of times each word appears in the same context as every other word.

Step—4: Use the co-occurrence matrix to compute the word embeddings using the GloVe algorithm. This involves training a model to minimize the error between the word vectors' dot product and the co-occurrence counts' logarithm.

Step—5: Save the resulting word embeddings to a file or use them directly in your model.

GloVe :Co-Occurrence Matrix

I love Programming. I love Math. I tolerate Biology.

	I	love	Program ming	Math	tolerate	Biology	.
I	0	2	0	0	1	0	2
love	2	0	1	1	0	0	0
Program ming	0	1	0	0	0	0	1
Math	0	1	0	0	0	0	1
tolerate	1	0	0	0	0	1	0
Biology	0	0	0	0	1	0	1
.	1	0	1	1	0	1	0

Window size = 1



The GloVe (Global Vector for word presentation)

Load the embeddings and use them as fixed word vectors in your application.

- <https://nlp.stanford.edu/projects/glove/>
- Download pre-trained word vectors
- Pre-trained word vectors. This data is made available under the [Public Domain Dedication and License](#) v1.0 whose full text can be found at: <http://www.opendatacommons.org/licenses/pddl/1.0/> [Wikipedia 2014](#) + [Gigaword 5](#) (6B tokens, 400K vocab, uncased, 50d, 100d, 200d, & 300d vectors, 822 MB download): [glove.6B.zip](#)
- Common Crawl (42B tokens, 1.9M vocab, uncased, 300d vectors, 1.75 GB download): [glove.42B.300d.zip](#)
- Common Crawl (840B tokens, 2.2M vocab, cased, 300d vectors, 2.03 GB download): [glove.840B.300d.zip](#)
- Twitter (2B tweets, 27B tokens, 1.2M vocab, uncased, 25d, 50d, 100d, & 200d vectors, 1.42 GB download): [glove.twitter.27B.zip](#)

glove.6B.300d , glove.6B.200d , glove.6B.100d , glove.6B.50d



Dealing with OOV

- Use a Default Vector
- Fallback to a Similar Word
- Train Your Own Embeddings

Better Solution: The FastText Model



The FastText Model

- The *FastText* model was introduced by **Facebook in 2016** as an extension and supposedly improvement of the vanilla Word2Vec model.
- FastText is a framework for learning word representations and performing robust, fast, and accurate text classifications

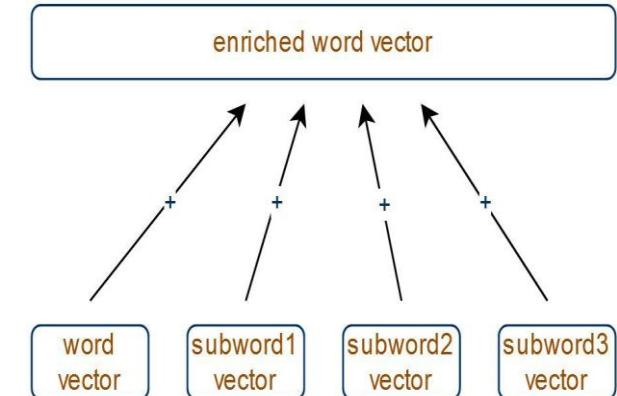
<https://fasttext.cc/>

FastText n-gram embedding model (Bojanowski et al., 2017): [Enriching Word Vectors with Subword Information](#)



The FastText Model...

Sub-word generation



Download English vector:

<https://dl.fbaipublicfiles.com/fasttext/vectors-crawl/cc.en.300.vec.gz>

Access from Kaggle: <https://www.kaggle.com/facebook/fasttext-wikinews>

FastText: Subword Generation

For a word, we generate character n-grams of length **3 to 6** present in it

Word	Length(n)	Character n-grams
eating	3	<ea, eat, ati, tin, ing, ng>
eating	4	<eat, eati, atin, ting, ing>
eating	5	<eati, eatin, ating, ting>
eating	6	<eatin, eating, ating>

Two-step vector representation updating

1. **First**, the embedding for the center word is calculated by taking a **sum of vectors for the character n-grams** and the whole word itself
2. For the actual context words, we directly take their word vector from the embedding table without adding the character n-grams



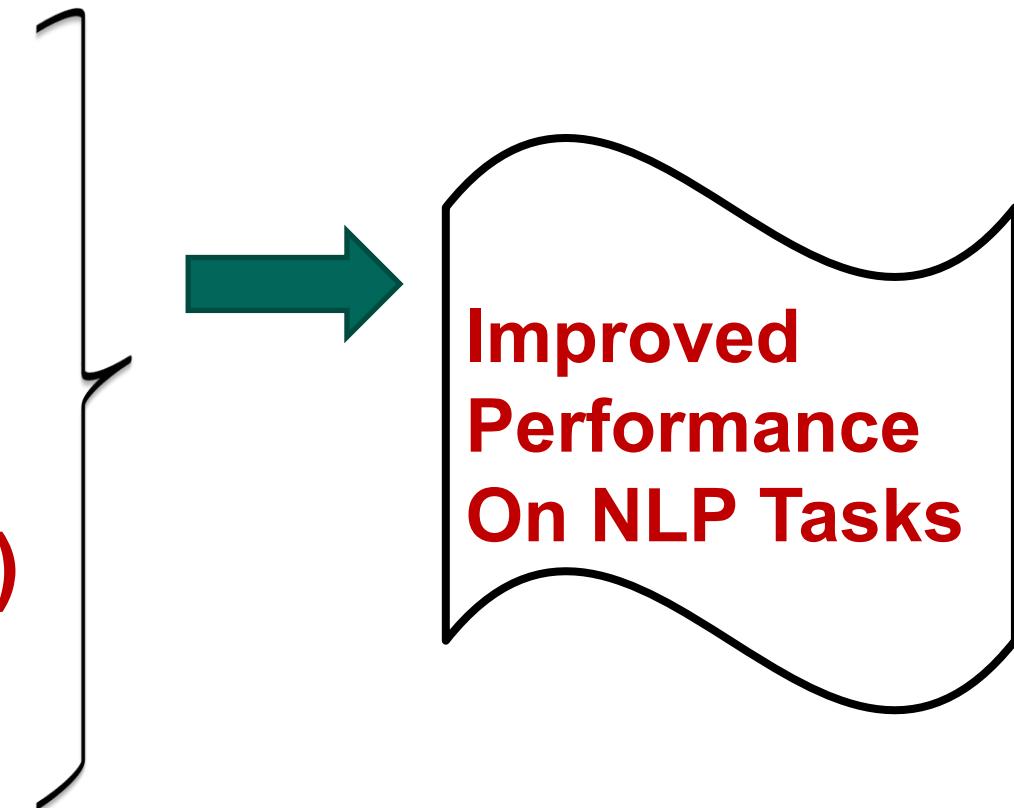
Advantages of FastText

- Capture fine level more gradual information
- Solve VOO
- Open-source, free, lightweight library.
- Handle text data in various languages .
- Has a simple and intuitive API



Word Embedding: Benefits

- Dimensionality reduction
- Semantic meaning
- Handling Out-of-Vocabulary (**OOV**)
- Transfer learning



Word Embedding - Limitations

Context Insensitivity

Bias

Limited Semantic Adaptation

Dimensionality

Resource Intensive

OOV words



Universal Text Representations

- *contextual word representations*
- Advanced neural language models
- complex architectures involving multiple passes through the text and multiple reads from left to right and right to left to model the context of language
 - **ELMo , BERT, ULMFiT**



Word Embedding - Evaluation

1. Intrinsic Evaluation

- Assessing the quality of word embeddings independently of any specific task.
- They focus on the internal properties of the embeddings. Word Similarity, Analogy Tasks, ...

1. Extrinsic Evaluation

- Assessing the quality of word embeddings based on their performance in downstream NLP tasks like: Text classification, NER, etc.



- ❖ WordNet and word senses
- ❖ Distributed representation
- ❖ Word Embedding
 - ❖ Word2Vec
 - ❖ Glove
 - ❖ FastText
- ❖ Evaluation of Word Embeddings
- ❖ Problems with Word Embedding



Q&A

