



CST8507: NATURAL LANGUAGE PROCESSING

WEEK#5
INTRODUCTION TO
LANGUAGE MODEL


DEVELOPED BY
HALA OWN, PH.D.

Lesson Agenda

- Lab 3
- Text Collection(overview)
- Language Model
 - N-gram
 - NN Language model
 - Recurrent Neural Networks RNN
 - LSTMs



NLP Development Life Cycle



Data generated in one minute on various social platforms




Image source: [HTTPs://localiq.com/blog/what-happens-in-an-internet-minute/](https://localiq.com/blog/what-happens-in-an-internet-minute/)


Text Collection

- Tweet Collecting
 - X API

The image shows a screenshot of the Twitter Developer Portal. On the left is a dark sidebar menu with white text and icons. The menu items are: 'Developer Portal' (with a Twitter icon), 'Dashboard' (with a house icon), 'Projects & Apps' (with a gear icon), 'Overview' (with a bar chart icon), 'Products' (with a product icon) and a green 'NEW' button, and 'Account' (with a gear icon). To the right of the sidebar is a light-colored dashboard area. At the top right of the dashboard is the word 'Dashboard'. Below it is a large icon of a laptop with a rocket launching from its screen. Underneath the icon is the text 'Create a Project to use v2 endpoints'. At the bottom center is a dark button with the text '+ Create Project' in white.



Create X Developer Account



Developer Portal Docs Community Updates Support

Twitter's v2 API
Build using all of Twitter's powerful v2 API endpoints

Write: 3,000 Tweets per month
Post up to 3,000 Tweets per month

Read: 10,000 Tweets per month
Retrieve up to 10,000 Tweets per month

If you need higher levels of access, [click here](#) to contact us and tell us more about your needs.

\$100.00 USD/month

[Subscribe to Basic](#)

[Sign up for Free Account](#)

A screenshot of the Twitter Developer Portal. At the top, there's a navigation bar with links for 'Developer Portal', 'Docs', 'Community', 'Updates', 'Support', and a user profile icon. Below the navigation, there's a large black box with the 'X' logo and the text 'Sign in to X'. To the right of this box, there's a section titled 'Twitter's v2 API' with a brief description. Below it are two sections: 'Write: 3,000 Tweets per month' and 'Read: 10,000 Tweets per month', each with a small icon and a brief description. Underneath these is a callout box with an info icon and text about higher access levels, followed by a note about the price (\$100.00 USD/month) and a 'Subscribe to Basic' button. At the bottom, there's a red-bordered 'Sign up for Free Account' button.

<https://help.rssground.com/articles/233141-how-to-create-x-twitter-developer-app>



ALGONQUIN
COLLEGE

Web Scraping: Extraction of data from a website

Python libraries are widely used for parsing HTML:

- 1. BeautifulSoup:** A popular library for parsing HTML and XML documents. It simplifies extracting data from web pages and has an active community with detailed documentation.
- 2. lxml:** Known for its speed, lxml is one of the fastest parsing libraries available. It receives regular updates, with the latest released in July 2023.
- 3. html5lib:** A pure-Python library designed to conform to the WHATWG (**Web Hypertext Application Technology Working Group**) HTML specification, ensuring compatibility with major web browsers.



Demo

- Inclass code



Reminder



PROBABILITY THEORY



ALGONQUIN
COLLEGE

Basic Probability Theory: Sampling with replacement

Pick a random shape, then put it back in the bag.



$$P(\square) = 2/15$$

$$P(\text{blue}) = 5/15$$

$$P(\text{blue} \mid \square) = 2/5$$

$$P(\square) = 1/15$$

$$P(\text{red}) = 5/15$$


$$P(\square) = 5/15$$

$$P(\square \text{ or } \triangle) = 2/15$$

$$P(\triangle \mid \text{red}) = 3/5$$

Sampling with replacement

Pick a random shape, then put it back in the bag.
What **sequence of shapes** will you draw?



$$\begin{aligned}P(\text{red, yellow, blue, blue}) &= 1/15 \times 1/15 \times 1/15 \times 2/15 \\&= 2/50625\end{aligned}$$

$$\begin{aligned}P(\text{blue, red, blue, red}) &= 3/15 \times 2/15 \times 2/15 \times 3/15 \\&= 36/50625\end{aligned}$$

$$P(\square) = 2/15$$

$$P(\text{blue}) = 5/15$$

$$P(\text{blue} | \square) = 2/5$$

$$P(\square) = 1/15$$

$$P(\text{red}) = 5/15$$

$$P(\square) = 5/15$$


$$P(\text{red or blue}) = 2/15$$

$$P(\triangle | \text{red}) = 3/5$$

Conditional Probability

$$P(X|Y) = \frac{P(X, Y)}{P(Y)}$$


The **conditional probability** of X given Y , Probability that one event occurs given that another event has already occurred.




$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$

is calculated as:

$$P(A | B) = P(B)$$



$P(B)$: Probability of B
All B Balls



$P(A | B)$: Probability of A given B
A within B (Purple Balls)

$$P(A \cap B) = P(A | B) \cdot P(B)$$

Chain Rule of Probability

The **chain rule** expresses a joint probability as a **product** of conditional probabilities.

For a sequence of events X_1, X_2, \dots, X_n

$$P(X_1, X_2, \dots, X_n) = P(X_1) P(X_2 | X_1) P(X_3 | X_1, X_2) \cdots P(X_n | X_1, \dots, X_{n-1})$$




LANGUAGE MODELING



Language Modeling:

the task of **predicting** what word comes next

the students opened their-----





Given a sequence of words $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(t)}, \mathbf{x}^{(t+1)}$
compute the probability distribution of the next word $P(\mathbf{x}^{(t+1)} | \mathbf{x}^{(t)}, \dots, \mathbf{x}^{(1)})$

Where $\mathbf{x}^{(t+1)}$ can be any word in the vocabulary $V = \{\mathbf{w}_1, \dots, \mathbf{w}_{|V|}\}$

- A system that does this is called a **Language Model**.

Popular Usages



A screenshot of a search results page titled "Phonetic spelling mistakes". It shows a search result for "Fizix is an interesting sudgetk." followed by a corrected version "Physics is an interesting subject." in a box.

Felix is an interesting subject 0.03
Felix is an interesting dialect 0.005
Physics is an interesting dialect 0.0034
Physics is an interesting subject 0.01

Goal of Language Modeling

learn patterns in text and predict the next word (or sequence of words) based on prior context.



N-gram Language Modeling

This is Big Data AI Book

<i>Uni-Gram</i>	This	Is	Big	Data	AI	Book
-----------------	------	----	-----	------	----	------

<i>Bi-Gram</i>	This is	Is Big	Big Data	Data AI	AI Book
----------------	---------	--------	----------	---------	---------

<i>Tri-Gram</i>	This is Big	Is Big Data	Big Data AI	Data AI Book
-----------------	-------------	-------------	-------------	--------------

IDEA: Collect statistics about how frequent different n-grams are, and use these to predict next word.

N-gram Language Modeling...

- For example, if we have sequence of tokens $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(T)}$ then the probability to see these tokens in this order is:

Using chain Rule

$$P(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(T)}) = P(\mathbf{x}^{(1)}) \times P(\mathbf{x}^{(2)} | \mathbf{x}^{(1)}) \times \dots \times P(\mathbf{x}^{(T)} | \mathbf{x}^{(T-1)}, \dots, \mathbf{x}^{(1)})$$

$$= \prod_{t=1}^T P(\mathbf{x}^{(t)} | \mathbf{x}^{(t-1)}, \dots, \mathbf{x}^{(1)})$$



This is what our LM provides



Language Modeling: n gram...

Our assumption $P(\mathbf{x}^{(t+1)} | \mathbf{x}^{(t)}, \dots, \mathbf{x}^{(1)}) = P(\mathbf{x}^{(t+1)} | \overbrace{\mathbf{x}^{(t)}, \dots, \mathbf{x}^{(t-n+2)}}^{n-1 \text{ words}})$

Recall the definition of conditional probabilities

$$p(B|A) = P(A, B) / P(A)$$
$$P(A, B) = P(A) P(B|A)$$

$\approx \frac{\text{count}(\mathbf{x}^{(t+1)}, \mathbf{x}^{(t)}, \dots, \mathbf{x}^{(t-n+2)})}{\text{count}(\mathbf{x}^{(t)}, \dots, \mathbf{x}^{(t-n+2)})}$

n-gram Language Models: Example using 4- gram

as the proctor started the clock the students opened their
discard fixed window

$$P(\mathbf{w}|\text{students opened their}) = \frac{\text{count(students opened their } \mathbf{w}\text{)}}{\text{count(students opened their)}}$$

For example, suppose that in the corpus:


- “students opened their” occurred 1000 times
 - “students opened their books” occurred 400 times
 - $\rightarrow P(\text{books} \mid \text{students opened their}) = 0.4$
 - “students opened their exams” occurred 100 times
 - $\rightarrow P(\text{exams} \mid \text{students opened their}) = 0.1$

N-grams : limitations and challenges


- Data Sparsity
- Computational Complexity
- Context Limitations



Neural Network Based Language Models




A Quick Review Of Neural Nets



- **Input layer** is a set of features; each arrow represents a weight (float number) that tells us how much each input contributes to each following step.
- Each node in the hidden layer is some combination of all the inputs. **The hidden layer** acts as the ‘input’ for the output layer.
- Backpropagation allows us to adjust the weights to improve accuracy and find the ‘correct’ way to combine the inputs and hidden layers to get the best possible results.

NN basic element: Perceptron or Neuron




Language Model: Neural Nets

as the proctor started the clock the students opened their -----?-----

discard

fixed window



Language Model: Neural Nets ...

output distribution

$$\hat{y} = \text{softmax}(U\mathbf{h} + \mathbf{b}_2) \in \mathbb{R}^{|V|}$$

hidden layer


$$\mathbf{h} = f(\mathbf{W}\mathbf{e} + \mathbf{b}_1)$$

concatenated word embeddings

$$\mathbf{e} = [\mathbf{e}^{(1)}; \mathbf{e}^{(2)}; \mathbf{e}^{(3)}; \mathbf{e}^{(4)}]$$

words / one-hot vectors

$$\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \mathbf{x}^{(3)}, \mathbf{x}^{(4)}$$



Feed Forward NN: Limitation...

as the proctor started the clock the students opened their
discard fixed window

Feed Forward NN: Limitation

“The food was good, not bad at all”

“The food was bad, not good at all”



Feed Forward NN: Limitation...

“Just watched the new movie. Loved it! 🍿 #entertained”

“The storyline was captivating, the characters were well-developed, and the cinematography was impressive. Overall, a fantastic movie night! 🎬 👍 #movienight #recommend”









Sequence Modeling: Motivations

- Handle **variable length** sequence data
- Track long **term dependency**
- Maintain information about **order**
- **Share information** across the sequence




DNN: Universal Approximation Theorem (UAT)



proven by George Cybenko in 1989

Core idea of Recurrent Neural Networks (RNNs) RNNs


Stateful computation



$$y_t, h_t = f(x_t, h_{t-1})$$


Core idea of RNNs ...

Stateful computation




$$h_t = W_x x_t + W_h h_{t-1} + b$$

How we train the model



These slides are sourced from Stanford's "Natural Language Processing with Deep Learning" course.


How we train the model




36

These slides are sourced from Stanford's "Natural Language Processing with Deep Learning" course.

How we train the model



How we train the model




38

These slides are sourced from Stanford's "Natural Language Processing with Deep Learning" course.

Loss $\longrightarrow J^{(1)}(\theta) + J^{(2)}(\theta) + J^{(3)}(\theta) + J^{(4)}(\theta) + \dots = J(\theta) = \frac{1}{T} \sum_{t=1}^T J^{(t)}(\theta)$

Predicted probability distribution



How we train the model

Language Model: RNN

output distribution $\hat{y}^{(4)} = P(x^{(5)}|\text{the students opened their})$

hidden states

$$h^{(t)} = \sigma(W_h h^{(t-1)} + W_e e^{(t)} + b_1)$$



$h^{(0)}$ is the initial hidden state

word embeddings


$$e^{(t)} = E x^{(t)}$$

$$x^{(t)} \in \mathbb{R}^{|V|}$$


$$\hat{y}^{(t)} = \text{softmax}(\mathbf{U} h^{(t)} + \mathbf{b}_2) \in \mathbb{R}^{|V|}$$



Difference between NN and RNN



Traditional NN for LM



RNN for LM


Image source: NLP in Action text book, O'Reilly

Fun With RNN Language Model

- <https://medium.com/@samim/obama-rnn-machine-generated-political-speeches-c8abd18a2ea0>




Back Propagation in RNN




Backpropagation Through Time (BPTT).

RNN Vanishing Gradient Intuition



These slides are sourced from Stanford's "Natural Language Processing with Deep Learning" course.


Vanishing gradient intuition



Vanishing gradient intuition


The Chain Rule

If $y = f(u)$, where $u = g(x)$


$$\frac{dy}{dx} = \frac{dy}{du} \cdot \frac{du}{dx}$$


These slides are sourced from Stanford's "Natural Language Processing with Deep Learning" course.

Vanishing gradient intuition



Vanishing gradient intuition



$$\frac{\partial J^{(4)}}{\partial h^{(1)}} = \frac{\partial h^{(2)}}{\partial h^{(1)}} \times$$

$$\frac{\partial h^{(3)}}{\partial h^{(2)}} \times$$


$$\frac{\partial h^{(4)}}{\partial h^{(3)}} \times \frac{\partial J^{(4)}}{\partial h^{(4)}}$$

chain rule!




These slides are sourced from Stanford's "Natural Language Processing with Deep Learning" course.

Vanishing gradient intuition



Why Vanishing Gradients is Problem

Vanishing gradients occur when the values of a gradient are too small and the model stops learning or takes way too long as a result



Vanishing Gradients Problem...

Example

*When she tried to print her **tickets**, she found that the printer was out of toner. She went to the stationery store to buy more toner. It was very overpriced. After installing the toner into the printer, she finally printed her-----*

RNN-LM needs to **model the dependency** between “tickets” on the 7th step and the target word “tickets” at the end



Long Short-Term Memory (LSTM)

- **Hochreiter & Schmidhuber (1997)** solved the problem of getting an RNN to remember things for a long time.
 - At each timestep t , the LSTM maintains two key components:
 - **Hidden state** – captures short-term dependencies.
 - **Cell state** – acts as a memory unit, storing long-term information.




Long Short-Term Memory (LSTM)

Key Concepts:

- Unlike standard RNNs, LSTMs can **control the flow of information** through three specialized gates:
 - **Forget gate** – decides which information to erase.
 - **Input gate** – determines what new information should be stored.
 - **Output gate** – regulates what information is passed to the next timestep.
- Each gate is represented as a vector of size n and can take values between **0 (closed)** and **1 (open)** dynamically, based on the current context.




Long Short -Term Memory (LSTM)



LSTM at time stamp T

Image source: <https://towardsdatascience.com/lstm-networks-a-detailed-explanation-8fae6aefc7f9>

Long Short Term Memory (LSTM): step 1




Forget gate: decide what parts of old state to forget

Long Short Term Memory (LSTM):step2

Input Gate:

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$


Input gate: decide how to update the cell state



Long Short Term Memory (LSTM):step3

Output Gate:


$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$



$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t * \tanh (C_t)$$

Finally, decide what to output as hidden state

Long Short Term Memory (LSTM)



LSTM Great resources

- <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>



Keras – Simplifying LSTMs in Python

Keras is a Python package that makes building and training TensorFlow neural networks really simple. We'll be working with the "Sequential" model which lets you add layers one at a time. As an example, let's see how to build a 1-layer LSTM model with 10 hidden nodes.

```
from keras.models import Sequential
from keras.layers import Dense, Activation, LSTM

model = Sequential()
model.add(LSTM(10, input_shape=(Timesteps, Feature_Length) ))
model.add(Dense(Number_of_Output_Nodes))
model.add(Activation('softmax'))
```



Evaluating Language Models

- The standard **evaluation metric** for Language Models is **perplexity**.

$$\text{perplexity} = \prod_{t=1}^T \left(\frac{1}{P_{\text{LM}}(\mathbf{x}^{(t+1)} | \mathbf{x}^{(t)}, \dots, \mathbf{x}^{(1)})} \right)^{1/T}$$

Inverse probability of corpus, according to Language Model

Normalized by
number of words

Low perplexity → the model predicts the text well

High perplexity → the text is unexpected for the model

Perplexity (PPL) measures **how confused** a language model is when predicting the next word in a sentence.



Summary

- We introduced the concepts of recurrent neural networks and how it can be applied to language problems.
- RNNs can be trained with a straightforward extension of the backpropagation algorithm.
- How LSTM used for text generation
- Applications of LSTM for sequence-to-sequence modeling



Q&A

