


OpenMMLab and CV Tech Stack

Stephin Rachel Thomas

June 30, 2025

What is OpenMMLab ?

1. An open-source tool system for computer vision
2. A big collection of state-of-the-art algorithms and dataset
3. A unified programming framework for efficient model development
4. A complete toolchain from model production to model deployment



The Philosophy Behind OpenMMLab

OpenMMLab is a comprehensive, open-source resource for computer vision research and development. It's built on the philosophy of **providing modular, reusable, and extendable components for various computer vision tasks**, ranging from object detection to action recognition. This approach simplifies the learning curve for researchers and developers, allowing them to focus on innovation rather than implementation details.

The Philosophy Behind OpenMMLab



Modular Approach in OpenMMLab

OpenMMLab adopts a modular approach, offering a suite of tools, each specialized for different computer vision tasks. This modular design allows users to select and combine components as needed, enhancing flexibility and efficiency. Toolboxes in OpenMMLab share a common framework, making it easier to switch between tasks or integrate multiple functionalities into a cohesive workflow.



Challenges of Diverse Source Codes and Models

One significant challenge in computer vision is the diversity of source codes and models available. Researchers often face difficulties in integrating and comparing different algorithms due to inconsistencies in implementation and documentation. OpenMMLab addresses this by providing standardized, well-documented codebases, enabling easier experimentation and comparison across various models and techniques.





OpenMMLab's Unified Interface for Computer Vision

OpenMMLab's unified interface across its toolboxes streamlines the process of developing and testing computer vision models. This consistency reduces the learning curve and development time, as users can apply similar methodologies and principles across different computer vision domains, be it segmentation, detection, or tracking.

MMPretrain: Pre-trained Model and Classification Toolbox Overview


MMPretrain, evolving within the OpenMMLab ecosystem, now encompasses not just a repository of pre-trained models but also focuses on image classification, integrating functionalities previously found in MMClassification. This expansion allows users to access state-of-the-art classification models and techniques, along with the robust pre-trained models for transfer learning, thereby catering to a broader range of computer vision tasks.



Out-of-Box Tasks


MMDetection: Object Detection Toolbox Explained

MMDetection is a versatile toolbox within OpenMMLab designed for object detection. It provides an extensive range of state-of-the-art detection algorithms, including Faster R-CNN, YOLO, and SSD. MMDetection is known for its high efficiency and flexibility, allowing researchers and developers to rapidly prototype and experiment with different detection models. Its modular design enables easy customization and extension, making it suitable for both academic research and industrial applications. Also has segmentation models.



MMDetection3D: 3D Object Detection Capabilities

MMDetection3D extends the capabilities of MMDetection to 3D object detection, catering to applications like autonomous driving and robotics. It supports various 3D detection frameworks, point cloud processing methods, and multi-modality fusion techniques. This toolbox simplifies working with 3D data, providing tools for 3D bounding box detection, point cloud segmentation, and LiDAR-camera fusion, thus enabling the development of sophisticated 3D perception models.



MMRotate: Focused on Rotation Detection

MMRotate is a specialized toolbox in OpenMMLab for handling rotation detection in images. It is particularly useful for aerial imagery, scene text detection, and other scenarios where objects are not aligned with the image axes. MMRotate includes various rotation-aware detection algorithms that can accurately detect and classify objects at arbitrary orientations, enhancing the performance of detection tasks in rotationally varied environments.



MMTracking: Video Object Tracking Features

MMTracking, another key component of OpenMMLab, focuses on video object tracking. It encompasses multiple algorithms for both single and multiple object tracking, accommodating different tracking scenarios from sports analytics to surveillance. MMTracking provides tools for real-time tracking, motion analysis, and trajectory prediction, making it a robust solution for dynamic and complex video sequences.





MMSegmentation: Semantic Segmentation Tools

MMSegmentation offers a comprehensive suite for segmentation tasks within the OpenMMLab framework. It includes a wide array of state-of-the-art segmentation models like U-Net, DeepLab, and PSPNet. This toolbox is designed for high performance and flexibility, supporting various segmentation scenarios such as medical image analysis, autonomous driving, and geographic information systems. MMSegmentation's modular design allows for easy experimentation and customization, facilitating the development of advanced segmentation models.



MMAction2: Action Recognition Toolbox Overview

MMAction2 is a comprehensive toolbox in OpenMMLab for action recognition and temporal action detection. It supports a wide range of action recognition models, including 3D CNNs and temporal segment networks. MMAction2 is suitable for applications in surveillance, human-computer interaction, and sports analysis, providing tools for analyzing and understanding complex actions and interactions in video data.



MMDeploy: Deployment Tools in OpenMMLab

MMDeploy is an open-source toolset designed for deploying deep learning models from the OpenMMLab ecosystem to various platforms and devices.

Model Converter:

Converts training models from OpenMMLab into backend models that can be run on target devices. Supports conversion to formats like ONNX, TorchScript, and others.

MMDeploy Model:


The result package exported by the Model Converter. Includes backend models and model meta information used by the Inference SDK

Inference SDK:

Developed in C/C++ and supports multiple languages such as Python, C#, and Java. Wraps preprocessing, model inference, and postprocessing modules.


Supported Platforms and Devices:

Compatible with various platforms including Linux, Windows, macOS, and Android. Supports multiple inference backends like ONNX Runtime, TensorRT, and OpenVINO. MMDeploy is particularly useful for deploying models in real-world applications, ensuring they run efficiently on different hardware setups.



The Importance of Conda in Computer Vision Tech Stacks

Conda is an essential tool for managing environments and dependencies in computer vision projects. It allows for the creation of isolated environments with specific versions of Python and libraries like TensorFlow and PyTorch, ensuring consistency and compatibility. Conda's environment management capabilities are particularly crucial for working with complex frameworks like OpenMMLab, helping to avoid dependency conflicts and streamline development workflows.



Understanding Conda: Basics and Installation


- Install libraries like NumPy, Pandas, Matplotlib for data handling and visualization: **conda install numpy pandas matplotlib**.
- Install OpenCV for image processing: **conda install -c conda-forge opencv**.
- For deep learning, install TensorFlow or PyTorch:
 - TensorFlow: **conda install -c conda-forge tensorflow**.
 - PyTorch: Visit the PyTorch website for the appropriate install command based on your system configuration.
- You can also use pip to install packages inside a conda environment
- To replicate the environment on another machine or share with others: `conda env export > environment.yml`.

Advanced Features of Conda for Dependency and Environment Management

- Conda excels in managing complex dependencies.
- Use 'conda list' to see installed packages.
- 'Conda env create -f environment.yml' creates an environment from a YAML file.
- 'Conda env list' shows all environments.
- Conda channels extend package availability.
- Resolve conflicts by specifying package versions

Introduction to Docker in Computer Vision

Docker offers portable, isolated environments for computer vision. It uses containers, lightweight and standalone executable packages. Containers run consistently across environments, ensuring that software runs the same everywhere. Docker simplifies the setup for complex computer vision projects, reducing 'works on my machine' issues and facilitating easier collaboration and deployment.




Leveraging Docker for Consistent Development and Deployment

Docker streamlines development and deployment. Create a Dockerfile to define the environment, then build it into an image using 'docker build'. Run this image as a container with 'docker run'. This process ensures that the development, testing, and production environments are identical. Docker containers can encapsulate OpenMMLab models, libraries, and dependencies, simplifying deployment and scaling.

```
1 ARG PYTORCH="1.9.0"
2 ARG CUDA="11.1"
3 ARG CUDNN="8"
4
5 FROM pytorch/pytorch:${PYTORCH}-cuda${CUDA}-cudnn${CUDNN}-devel
6
7 ENV TORCH_CUDA_ARCH_LIST="6.0 6.1 7.0 7.5 8.0 8.6+PTX" \
8     TORCH_NVCC_FLAGS="-Xfatbin -compress-all" \
9     CMAKE_PREFIX_PATH="$(dirname $(which conda))/.." \
10    FORCE_CUDA="1"
11
12 # Avoid Public GPG key error
13 # https://github.com/NVIDIA/nvidia-docker/issues/1631
14 RUN rm /etc/apt/sources.list.d/cuda.list \
15     && rm /etc/apt/sources.list.d/nvidia-ml.list \
16     && apt-key del 7fa2af80 \
17     && apt-key adv --fetch-keys https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1804/x86_64/3bf863cc.pub \
18     && apt-key adv --fetch-keys https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1804/x86_64/7fa2af80.pub
19
20 # (Optional, use Mirror to speed up downloads)
21 # RUN sed -i 's/http://\//archive.ubuntu.com/ubuntu\//http://\//mirrors.aliyun.com/ubuntu\//g' /etc/apt/sources.list && \
22 #     pip config set global.index-url https://pypi.tuna.tsinghua.edu.cn/simple
23
24 # Install the required packages
25 RUN apt-get update \
26     && apt-get install -y ffmpeg libsm6 libxext6 git ninja-build libglib2.0-0 libsm6 libxrender-dev libxext6 \
27     && apt-get clean \
28     && rm -rf /var/lib/apt/lists/*
29
30 # Install MMEngine and MMCV
31 RUN pip install openmim && \
32     mim install "mmengine>=0.7.1" "mimcv>=2.0.0rc4"
33
34 # Install MMDetection
35 RUN conda clean --all \
36     && git clone https://github.com/open-mmlab/mmdetection.git /mmdetection \
37     && cd /mmdetection \
38     && pip install --no-cache-dir -e .
39
40 WORKDIR /mmdetection
```


Docker Hub: Open Source Images and Repository for Your Projects

Docker Hub is a cloud-based repository for managing Docker images. It hosts numerous open-source images, which can be used as the basis for custom containers. Users can pull images from Docker Hub using 'docker pull' and push their own images with 'docker push'. It facilitates sharing and collaboration, allowing teams to easily distribute and manage Docker images. Docker Hub also supports private repositories for confidential projects.



Visual Studio Code for Computer Vision Development

Visual Studio Code (VS Code) is a versatile code editor for computer vision development. It supports Python and other programming languages with features like IntelliSense for code completion and debugging tools. Extensions like Python, Docker, and Git enhance functionality. VS Code's integration with version control systems and its lightweight design make it ideal for developing complex computer vision projects.



VS Code Features: Plugins, Remote Debugging, and More

VS Code offers a wide range of features and plugins that enhance productivity. The Python extension supports linting, testing, and environment management. The Live Share extension enables real-time collaborative coding. Remote Development plugins allow coding on remote systems like Docker containers or cloud servers. VS Code's debugging tools, including breakpoints, call stack inspection, and variable exploration, simplify problem-solving in complex codebases.



Debugging in Docker with VS Code

VS Code can debug applications running inside Docker containers. By using the Remote - Containers extension, developers can attach to a running container and debug using VS Code's powerful debugging tools. This setup allows for testing in an environment identical to production. It simplifies the process of diagnosing and fixing issues in containerized computer vision applications, ensuring consistency across development and deployment stages.

Utilizing AWS EC2 for Compute-Intensive Tasks in Computer Vision

Amazon Web Services (AWS) EC2 provides scalable compute capacity in the cloud, ideal for compute-intensive computer vision tasks. EC2 offers a wide range of instance types, including GPU-enabled instances for deep learning tasks. Users can easily scale their compute resources up or down based on demand, making EC2 a flexible and cost-effective solution for training models, processing large datasets, and deploying computer vision applications.



Optimizing Computer Vision Workloads on AWS EC2

Optimizing workloads on AWS EC2 involves selecting the right instance types, managing storage efficiently, and leveraging AWS's networking capabilities. For deep learning, choosing GPU instances like the P3 or G4 series can significantly speed up model training. Efficient use of Elastic Block Store (EBS) and Amazon S3 for data storage and management is crucial. Additionally, using AWS's networking features can improve data transfer speeds and reduce latency, enhancing the overall performance of computer vision applications.




Advanced Use of AWS EC2: Scalability and Cost Management

AWS EC2 excels in scalable computing, allowing users to adjust resources as per project demands. It's vital for handling varying workloads, especially in large-scale computer vision projects. Utilize Auto Scaling to adjust capacity and maintain performance. Cost management tools like AWS Budgets and Cost Explorer help monitor and optimize expenses. Spot Instances offer cost savings for flexible workloads.


Community and Support in OpenMMLab Ecosystem

The OpenMMLab ecosystem is supported by a vibrant community of developers and researchers. Users can access extensive documentation, tutorials, and GitHub repositories for each toolbox. Community forums and platforms like Stack Overflow offer support and discussion opportunities. Regular updates and contributions from users around the world keep the toolboxes state-of-the-art and user-friendly.



TensorFlow Object Detection API: Overview

The TensorFlow Object Detection API is a powerful toolkit for building object detection models. It provides pre-trained models, multiple architectures, and the ability to train custom detectors. This API has been instrumental in advancing object detection research and applications. It supports various models like SSD, Faster R-CNN, and Mask R-CNN, making it a versatile tool for different detection tasks.



Conclusion: Leveraging OpenMMLab for Advanced Computer Vision

In conclusion, OpenMMLab provides a comprehensive suite of tools that cater to various aspects of computer vision. Its modular design, coupled with support from the community, makes it an ideal choice for both academic research and industrial applications. By leveraging OpenMMLab in conjunction with tools like Conda, Docker, VS Code, and AWS EC2, developers and researchers can build, deploy, and scale advanced computer vision models more efficiently and effectively.