


# Gymnasium Environments

# Agenda

- This week:
  - Gymnasium custom environment: an upgrade for our Lab 1 CliffWalking environment class
  - Pygame for rendering
  - BlocksWorld-v0 environment (future Assignment 1)
  - stable-baselines3

# What we've implemented so far

- Lab 1 Agent: <your\_algonquin\_id>\_lab2\_qlearning\_agent.py
- Lab 1 Environment: <your\_algonquin\_id>\_lab2\_cliff\_env.py (simple "homemade" gridworld and CliffWalking)



**Figure 3.1:** The agent–environment interaction in a Markov decision process.

- Lab 2: <your\_algonquin\_id>\_lab2\_cliff\_env.py replaced with gymnasium environment
- Lab 2: <your\_algonquin\_id>\_lab2\_qlearning\_agent.py updated to work with gymnasium environment

# Where the implementation is going

- Lab 2: upgrade our cliffwalking environment to Gymnasium
- Turn the Prolog blocks world into a Gymnasium BlocksWorld-v0 environment (assignment 1)
- Use PyGame for rendering our environments
- Our <your\_algonquin\_id>\_lab2\_qlearning\_agent.py implementation still plays the role of agent
- stable-baselines3 is a set of standard RL algorithms that work with Gymnasium environments
- We can use the stable-baselines3 algorithms: DQN, PPO, for new agents



# Diagram comparison (Assignment 1)






Figure 3.1: The agent–environment interaction in a Markov decision process.



# Important Sites

- Basic Components from Lab1:

<https://medium.com/data-science/math-of-q-learning-python-code-5dcbdc49b6f6>

Custom environment for Gymnasium Lab 2:

[https://gymnasium.farama.org/tutorials/gymnasium\\_basics/environment\\_creation/#](https://gymnasium.farama.org/tutorials/gymnasium_basics/environment_creation/#)

- PyGame is often used by Gymnasium Environments for rendering:

[https://dr0id.bitbucket.io/legacy/pygame\\_tutorial00.html](https://dr0id.bitbucket.io/legacy/pygame_tutorial00.html)

[https://dr0id.bitbucket.io/legacy/pygame\\_tutorial01.html](https://dr0id.bitbucket.io/legacy/pygame_tutorial01.html)

- Stable-baselines3:

<https://stable-baselines3.readthedocs.io/en/master/guide/install.html>



# Gym to Gymnasium

- Gymnasium is common as the successor to Gym

```
import gymnasium as gym
```

- Reset method takes seed as parameter (for reproducing random sequences)

```
reset(self, seed=None, options=None)
```

- Step method returns truncated

```
return observation, reward, terminated, truncated, info
```

# Creating a Gymnasium Environment

[https://gymnasium.farama.org/tutorials/gymnasium\\_basics/environment\\_creation/#](https://gymnasium.farama.org/tutorials/gymnasium_basics/environment_creation/#)

Steps to creating a new Gymnasium environment:

- import gymnasium as gym
- from gymnasium import spaces # note: from gymnasium, not from gym
- Make the <somename>Env python class extend gymnasium.env
- Implement/override the gymnasium.env methods:
  - `__init__(self, render_mode=None, size=5)`
  - `reset(self, seed=None, options=None)`
  - `step(self, action)->return observation, reward, terminated, False, info`
  - `render(self)`
  - `close(self)`



# Creating a Gymnasium Environment

- Define the state space and action space
- Example from gymnasium docs:

# Observations for the gridworld example are dictionaries with the agent's and the target's location.

# Each location is encoded as an element of {0, ..., `size`}^2, i.e. MultiDiscrete([size, size]).

```
self.observation_space = spaces.Dict(
```

```
{
```

```
    "agent": spaces.Box(0, size - 1, shape=(2,), dtype=int),
```

```
    "target": spaces.Box(0, size - 1, shape=(2,), dtype=int),
```

```
)
```

# We have 4 actions, corresponding to "right", "up", "left", "down"

```
self.action_space = spaces.Discrete(4)
```

# Creating a Gymnasium Environment

- Define the state space and action space
- In this case, the observation is a dictionary of 3-letter-state integers:

```
self.observation_space = spaces.Dict(  
    {  
        "agent": spaces.Discrete(len(self.states)),  
        "target": spaces.Discrete(len(self.states)),  
    }  
)  
  
self.action_space = spaces.Discrete(len(self.actions))
```

# Creating a Gymnasium Environment

- Define the state space and action space
- In this case (per the coming Assignment 1 document), the observation is a single larger integer representing both the current state and target at the same time:

```
self.observation_space = spaces.Discrete(len(self.states))
```

```
self.action_space = spaces.Discrete(len(self.actions))
```

# StableBaselines MlpPolicy vs MultiInputPolicy

- Dictionary observations:

```
self.observation_space = spaces.Dict(  
    {  
        "agent": spaces.Discrete(len(self.states)),  
        "target": spaces.Discrete(len(self.states)),  
    }  
)  
  
self.action_space = spaces.Discrete(len(self.actions))
```

- In the agent

```
# For spaces.Dict observations, use MultiInputPolicy  
  
import gymnasium as gym  
  
env = gym.make("aisd_examples/BlocksWorld-v0", render_mode="human")  
  
model = DQN("MultiInputPolicy", env, verbose=1)  
  
model.learn(total_timesteps=10000, log_interval=4)
```



# StableBaselines MlpPolicy vs MultilnputPolicy

- Observation is a discrete number:

```
self.observation_space = spaces.Discrete(len(self.states))
```

```
self.action_space = spaces.Discrete(len(self.actions))
```

- In the agent

```
# for spaces.Discrete observation_space, use MlpPolicy (Multi-Layer-Perceptron)
```

```
import gymnasium as gym
```

```
env = gym.make("aisd_examples/BlocksTarget-v0", render_mode="human")
```

```
model = DQN("MlpPolicy", env, verbose=1)
```

```
model.learn(total_timesteps=10000, log_interval=4)
```

# Put our environment in a package to install

As outlined in the Gymnasium docs (they use the name gymnasium\_env):

```
<argonquin_id>_blocksworld_env/  
    pyproject.toml  
    blocksworld_env/      # gymnasium_env/ would be at this level  
        __init__.py  
    envs/  
        __init__.py  
        blocks_world.py
```



# Using our custom environment

Using our custom blocksworld environment with Discrete observations

<your\_algonquin\_id>\_lab2\_qlearning\_agent.py from Lab 2:

```
#from medium_qlearning_env import Env  
  
import blocksworld_env  
  
import gymnasium as gym  
  
import matplotlib.pyplot as plt  
  
import numpy as np  
  
import time  
  
import os  
  
# create environment  
  
#env = Env()  
  
env = gym.make("blocksworld_env/BlocksWorld-v0", render_mode="human")  
  
# QTable : contains the Q-Values for every (state,action) pair  
  
qtable = np.random.rand(env.observation_space.n, env.action_space.n).tolist()
```

# Using our custom environment

Using our environment (Discrete observations) with stable-baselines3:

```
import gymnasium as gym

import blocksworld_env

from stable_baselines3 import DQN

env = gym.make("blocks_world/BlocksWorld-v0", render_mode="human")

model = DQN("MlpPolicy", env, verbose=1)          # for non-dictionary Discrete observation where state and target are combined

model.learn(total_timesteps=10000, log_interval=4)

model.save("dqn_blocks")

del model # remove to demonstrate saving and loading

model = DQN.load("dqn_blocks")

obs, info = env.reset()

while True:

    action, _states = model.predict(obs, deterministic=True)

    obs, reward, terminated, truncated, info = env.step(action)

    if terminated or truncated:

        obs, info = env.reset()
```

# Prolog Blocks world environment (Assignment 1)

```
% situation calculus blocks world (just one fluent, clear is defined with on)
:- dynamic on/3.

% target means this is the goal (not used because Python picks random state)
target:-on(a,4,[],on(b,c,[],on(c,a,[]).

% reset asserts the initial state
reset:-
    retractall(on(_,_,[])),
    assert(on(a,1,[])),
    assert(on(b,3,[])),
    assert(on(c,a,[])).
```

# Prolog Blocks world environment

```
state(S) :-  
    (block(A); place(A)), dif(A,a),  
    (block(B); place(B)), dif(B,b),  
    (block(C); place(C)), dif(C,c),  
    dif(A,B),  
    dif(B,C),  
    dif(A,C),  
    grounded(A,B,C),  
    atomics_to_string([A,B,C],S).  
  
grounded(A,B,C) :-  
    legal(A,B,C),  
    (place(A); place(B); place(C)), !.  
  
legal(A,B,C) :-  
    block(A), block(B), A=b, B=a, !, fail  
;  
    block(B), block(C), B=c, C=b, !, fail  
;  
    block(A), block(C), A=c, C=a, !, fail  
;  
    true.
```

# Prolog Blocks world environment

```
% action(Act) means Act is conceivable  
  
action(Act):-  
    Act = move(A,B,C),  
    block(A),  
    (block(B);place(B)),  
    (block(C);place(C)),  
    dif(A,B),  
    dif(A,C),  
    dif(B,C).
```

```
current_state(S):-  
    on(a,A,[ ]),  
    on(b,B,[ ]),  
    on(c,C,[ ]),  
    atomics_to_string([A,B,C],S).  
  
step(Act):- poss([Act]),  
    on(a,A,[Act]),  
    on(b,B,[Act]),  
    on(c,C,[Act]),  
    retractall(on(_,_,[ ])),  
    assert(on(a,A,[ ])),  
    assert(on(b,B,[ ])),  
    assert(on(c,C,[ ])).
```

# PyGame for rendering

```
# import the pygame module, so you can use it
import pygame

class Display():
    def __init__(self):
        # initialize the pygame module
        pygame.init()

        # load and set the logo
        logo = pygame.image.load("logo32x32.png")
        pygame.display.set_icon(logo)
        pygame.display.set_caption("Blocks World")

        # create a surface on screen that has the given size
        # define a variable to control the main loop
        self.running = True

        IMAGE_SIZE_X = 100
        IMAGE_SIZE_Y = 100
        DEFAULT_IMAGE_SIZE = (IMAGE_SIZE_X, IMAGE_SIZE_Y)
        self.screen = pygame.display.set_mode((4*IMAGE_SIZE_X,6*IMAGE_SIZE_Y))

#see screen.py on brightspace
```

